

# CSE 141 – Computer Architecture

## Summer Session I, 2004

### Lectures 8

### Pipelining

Pramod V. Argade

# CSE141: Introduction to Computer Architecture

**Instructor:** Pramod V. Argade (p2argade@cs.ucsd.edu)

Office Hours:

Tue. 7:30 - 8:30 PM (AP&M 4141)

Wed. 4:30 - 5:30 PM (AP&M 4141)

**TA:**

Anjum Gupta (a3gupta@cs.ucsd.edu)

Office Hour: Mon/Wed 12 - 1 PM

Chengmo Yang (c5yang@cs.ucsd.edu)

Office Hour: Mon/Thu 2 - 3 PM

**Lecture:** Mon/Wed. 6 - 8:50 PM, Center 109

**Textbook:** Computer Organization & Design  
The Hardware Software Interface, 2<sup>nd</sup> Edition.  
Authors: Patterson and Hennessy

**Web-page:** <http://www.cse.ucsd.edu/classes/su04/cse141>

# Announcements

- Reading Assignment:
  - Pipelining, Sections 6.1 - 6.3 (Tuesday)
  - Pipelining, Sections 6.4 - 6.7 (Wednesday)
- Homework 5: Due Mon., July 26 in class
  - 5.29
  - 6.4, 6.10, 6.11, 6.12, 6.13, 6.20, 6.23, 6.26
- Quiz
  - When:** Mon, July 26, First 10 minutes of the class
  - Topic:** Pipeline Hazards, Chapter 6
  - Need:** Paper, pen
- Final Exam
  - When:** Sat., July 31, 7 - 10 PM, Center 109 (Time and Room may change!)

# CSE141 Course Schedule

Lecture #	Date	Time	Room	Topic	Quiz topic	Homework Due
1	Mon. 6/28	6 - 8:50 PM	Center 109	Introduction, Ch. 1 ISA, Ch. 3	-	-
2	Wed. 6/30	6 - 8:50 PM	Center 109	Performance, Ch. 2 Arithmetic, Ch. 4	ISA Ch. 3	#1
-	Mon. 7/5	<b>No Class</b>		July 4th Holiday	-	-
3	Wed. 7/7	6 - 8:50 PM	Center 109	Arithmetic, Ch. 4 Cont. Single-cycle CPU Ch. 5	Performance Ch. 2	#2
4	Mon. 7/12	6 - 8:50 PM	Center 109	Single-cycle CPU Ch. 5 Cont. Multi-cycle CPU Ch. 5	Arithmetic, Ch. 4	#3
5	Tue. 7/13	7:30 - 8:50 PM	Center 109	Multi-cycle CPU Ch. 5 Cont. <b>(July 5th make up class)</b>	-	-
6	Wed. 7/14	6 - 8:50 PM	Center 109	Single and Multicycle CPU Examples and Review for Midterm	Single-cycle CPU Ch. 5	-
7	Mon. 7/19	6 - 8:50 PM	Center 109	Mid-term Exam Exceptions	-	#4
8	Tue. 7/20	7:30 - 8:50 PM	Center 109	Pipelining Ch. 6 <b>(July 5th make up class)</b>	-	-
9	Wed. 7/21	6 - 8:50 PM	Center 109	Hazards, Ch. 6	-	-
10	Mon. 7/26	6 - 8:50 PM	Center 109	Memory Hierarchy & Caches Ch. 7	Hazards Ch. 6	#5
11	Wed. 7/28	6 - 8:50 PM	Center 109	Virtual Memory, Ch. 7 Course Review	Cache Ch. 7	#6
12	Sat. 7/31	7 - 10 PM	Center 109	Final Exam	-	-

# Would our pipeline design work in any case?

- What happens when...

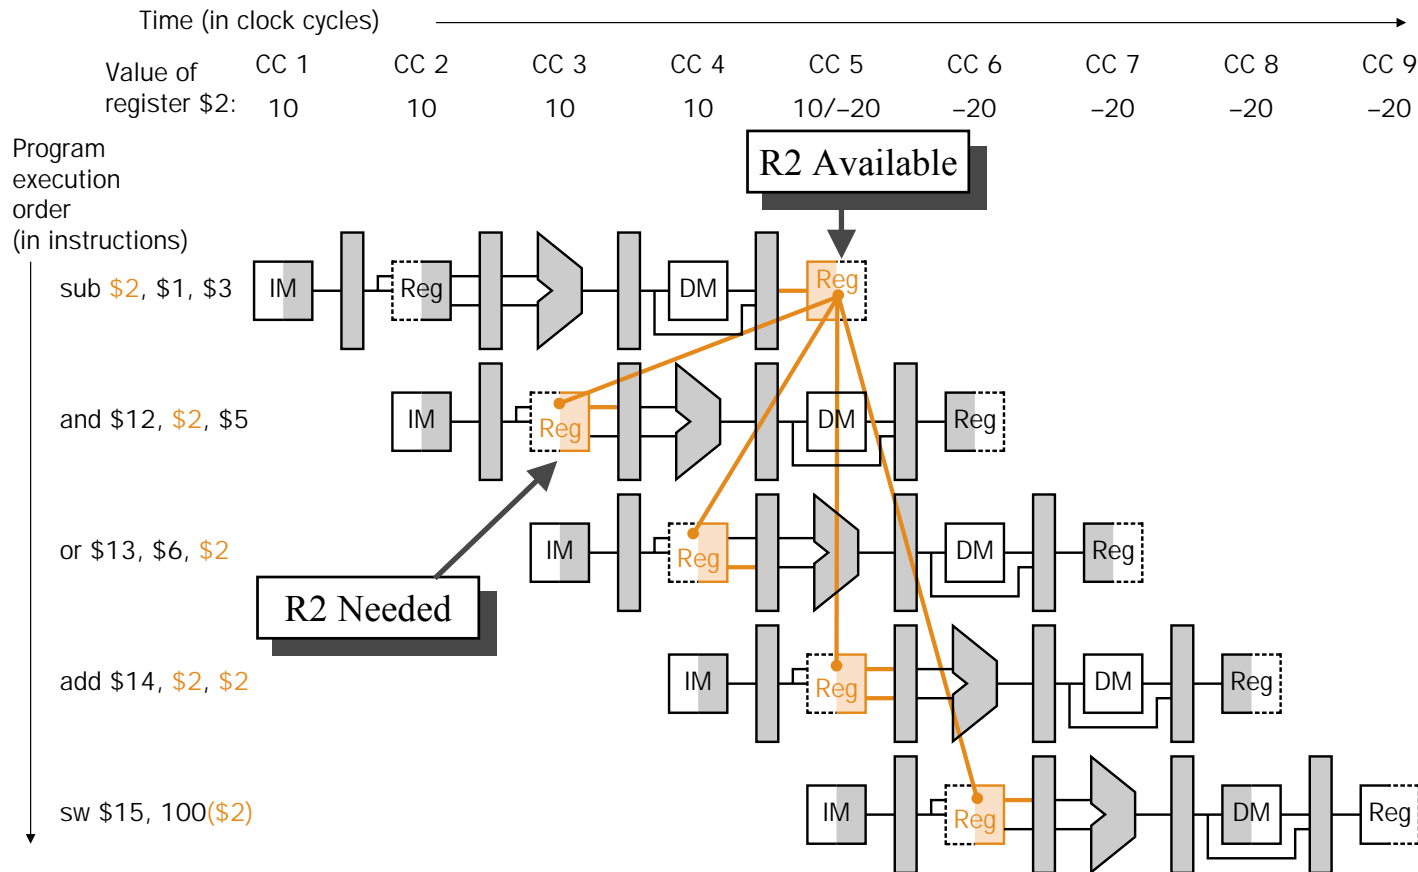
add \$3, \$10, \$11

lw \$8, 1000(\$3)

sub \$11, \$8, \$7

# Data Hazards

- When a result is needed in the pipeline before it is available, a “data hazard” occurs.



- Result of SUB instruction not available until CC5 or later!**

# Software Solutions to Data Hazards

- Have compiler guarantee no hazards
  - Rearrange code to remove hazard
    - Not possible every time

- Insert “nops”

- Where do we insert the “nops” ?

```
sub    $2, $1, $3
and    $12, $2, $5
or     $13, $6, $2
add    $14, $2, $2
sw     $15, 100($2)
```

nop  
nop

- Problem: Data hazards are very common!
  - “nops” really slows us down!

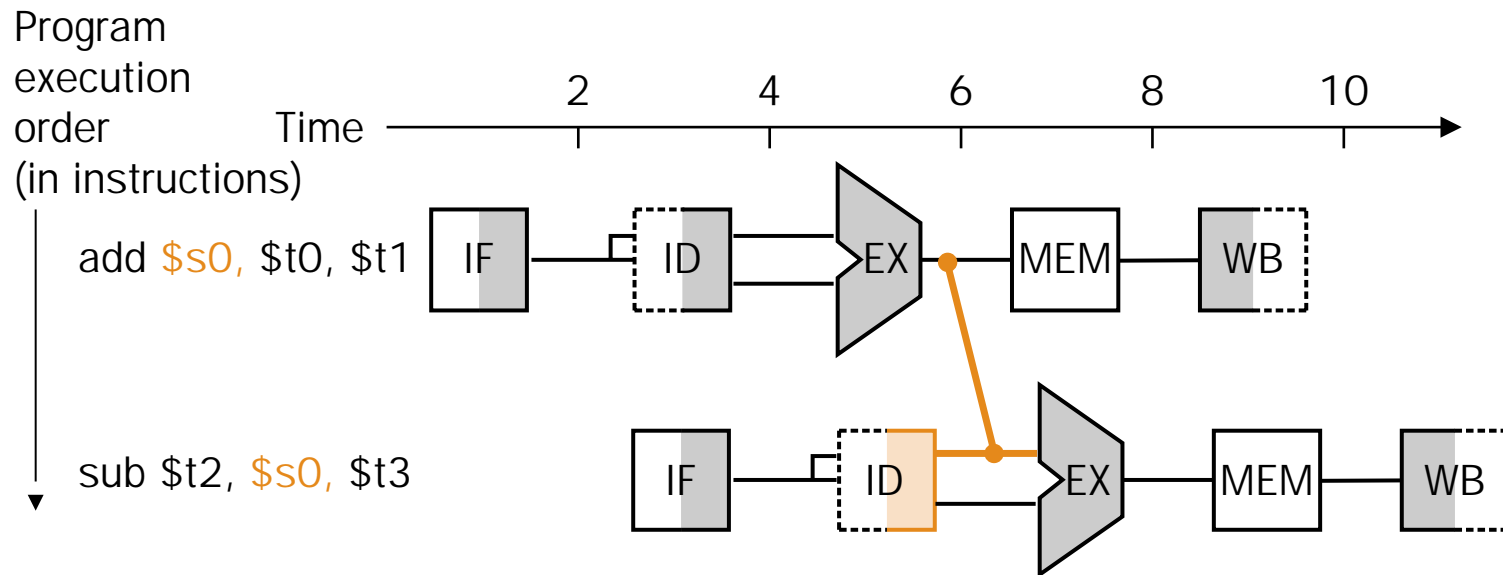
# Hardware Solutions to Data Hazards

- Stall the pipeline (insert bubbles)
  - Data hazards are too common
    - Same as “nops”
  - Severe performance hit
- Forward the data as soon as it is available
  - Modify the pipeline to forward (bypass data)

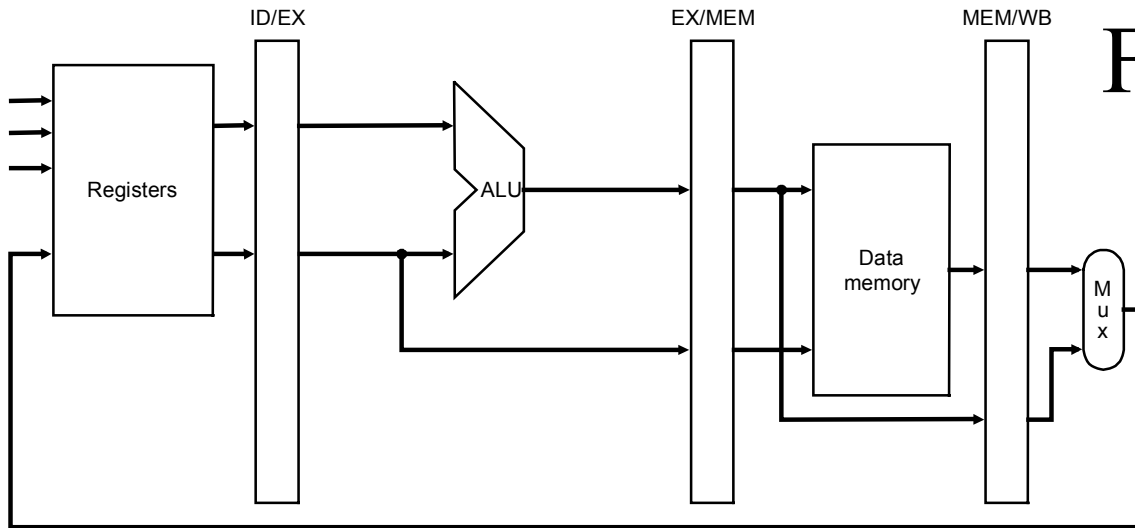


# Forwarding

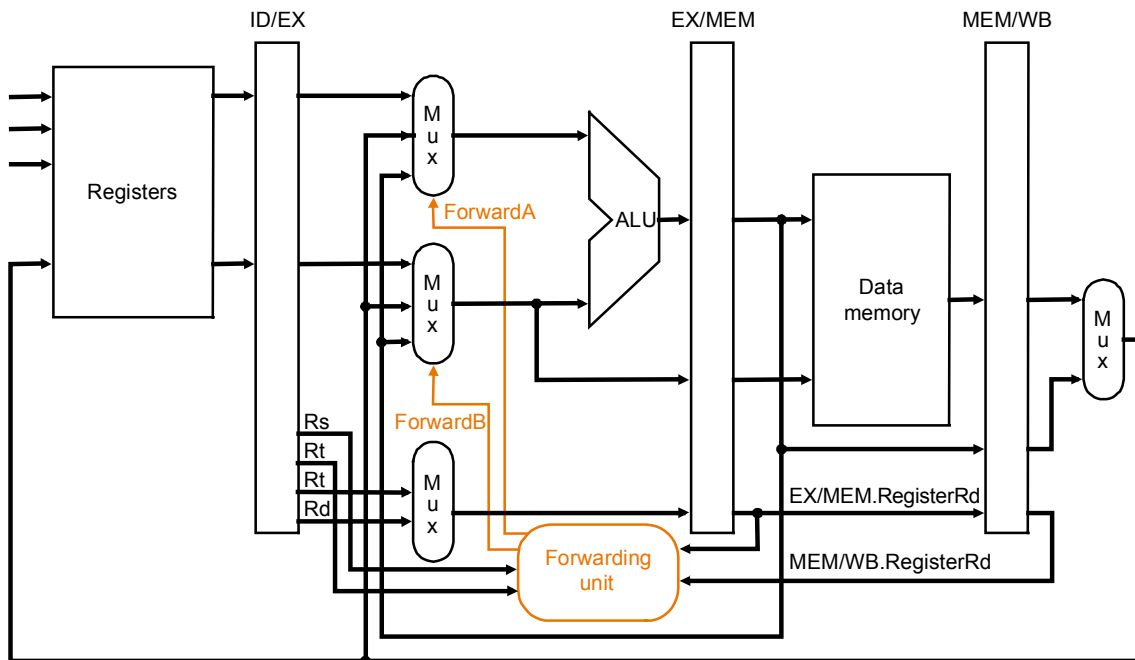
- Use temporary results, don't wait for them to be written



# Forwarding

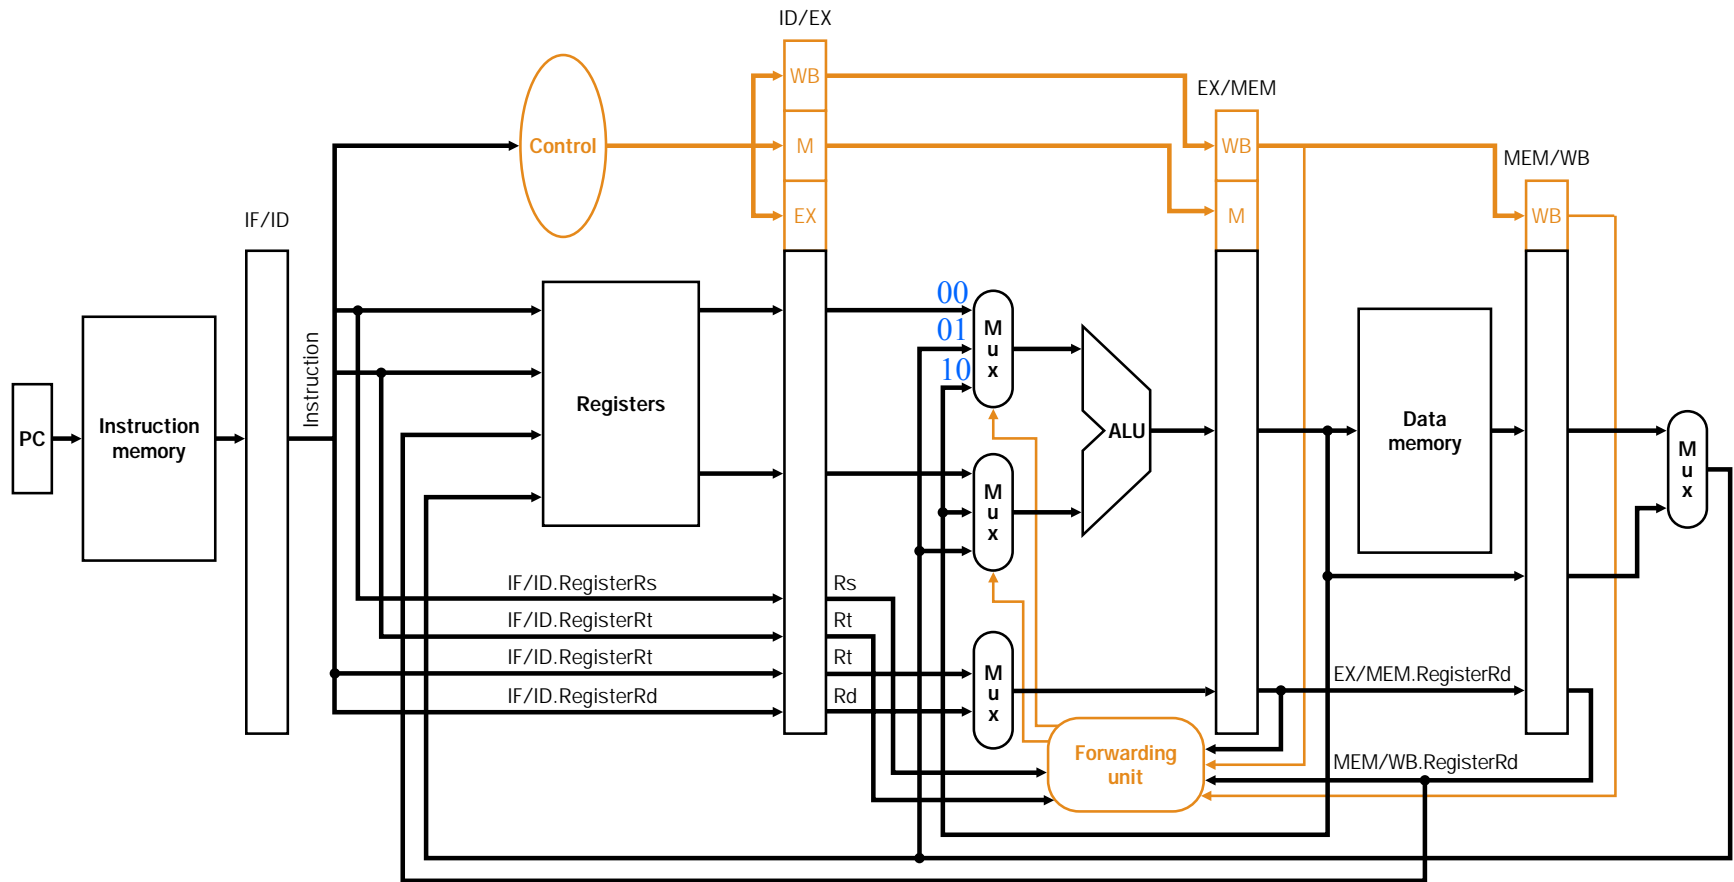


a. No forwarding




b. With forwarding

# Reducing EX Data Hazards Through Forwarding



if (EX/MEM.RegWrite  
 and (EX/MEM.RegisterRd != 0)  
 and (EX/MEM.RegisterRd = ID/EX.RegisterRs)) ForwardA = 10  
 if (EX/MEM.RegWrite  
 and (EX/MEM.RegisterRd != 0)  
 and (EX/MEM.RegisterRd = ID/EX.RegisterRt)) ForwardB = 10

# Reducing MEM Data Hazards Through Forwarding



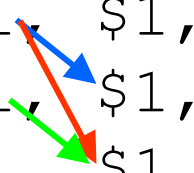
```

if (MEM/WB.RegWrite
    and (MEM/WB.RegisterRd != 0)
    and (MEM/WB.RegisterRd = ID/EX.RegisterRs)) ForwardA = 01
if (MEM/WB.RegWrite
    and (MEM/WB.RegisterRd != 0)
    and (MEM/WB.RegisterRd = ID/EX.RegisterRt)) ForwardB = 01
    
```

# Simultaneous EX/MEM Forwarding

- Consider following code

```
add $1, $1, $2  
add $1, $1, $3  
add $1, $1, $4
```



...

- Must forward from MEM stage
- Disable WB stage forwarding

if (MEM/WB.RegWrite

and (MEM/WB.RegisterRd != 0)

**and (EX/MEM.RegisterRd ≠ ID/EX.RegisterRs)**

and (MEM/WB.RegisterRd = ID/EX.RegisterRs)) ForwardA = 01

if (MEM/WB.RegWrite

and (MEM/WB.RegisterRd != 0)

**and (EX/MEM.RegisterRd ≠ ID/EX.RegisterRt)**

and (MEM/WB.RegisterRd = ID/EX.RegisterRt)) ForwardB = 01

# Forwarding in Action

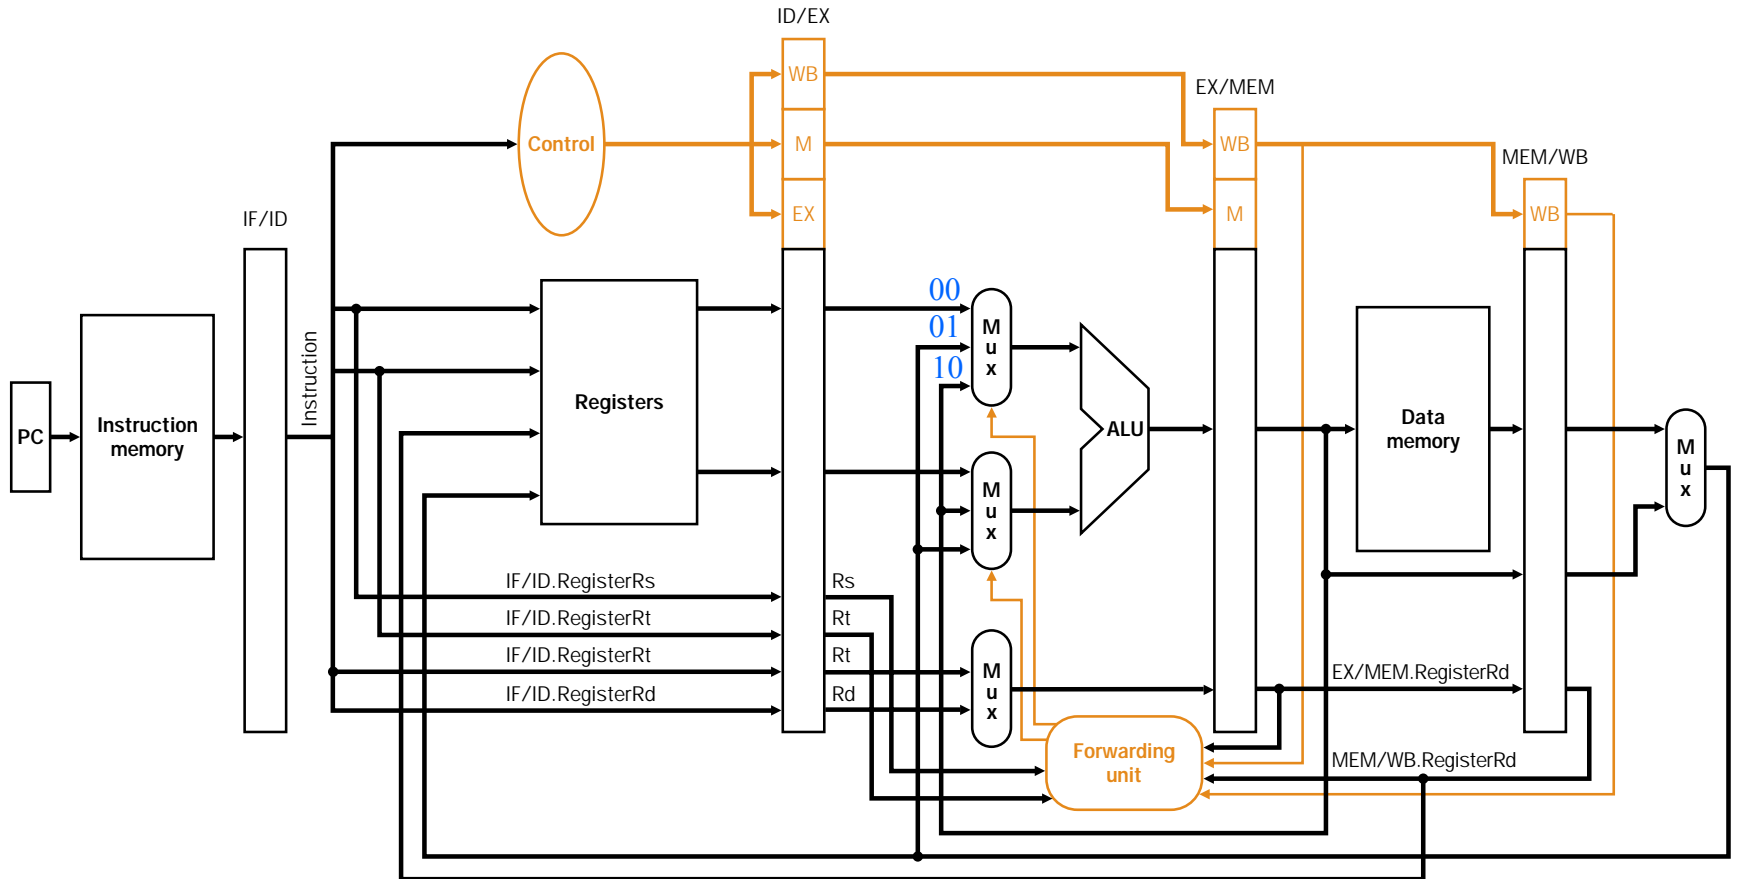
sub \$1, \$12, \$3

and \$12, \$3, \$4

add \$3, \$8, \$11

Memory Access

Write Back



# Forwarding in Action

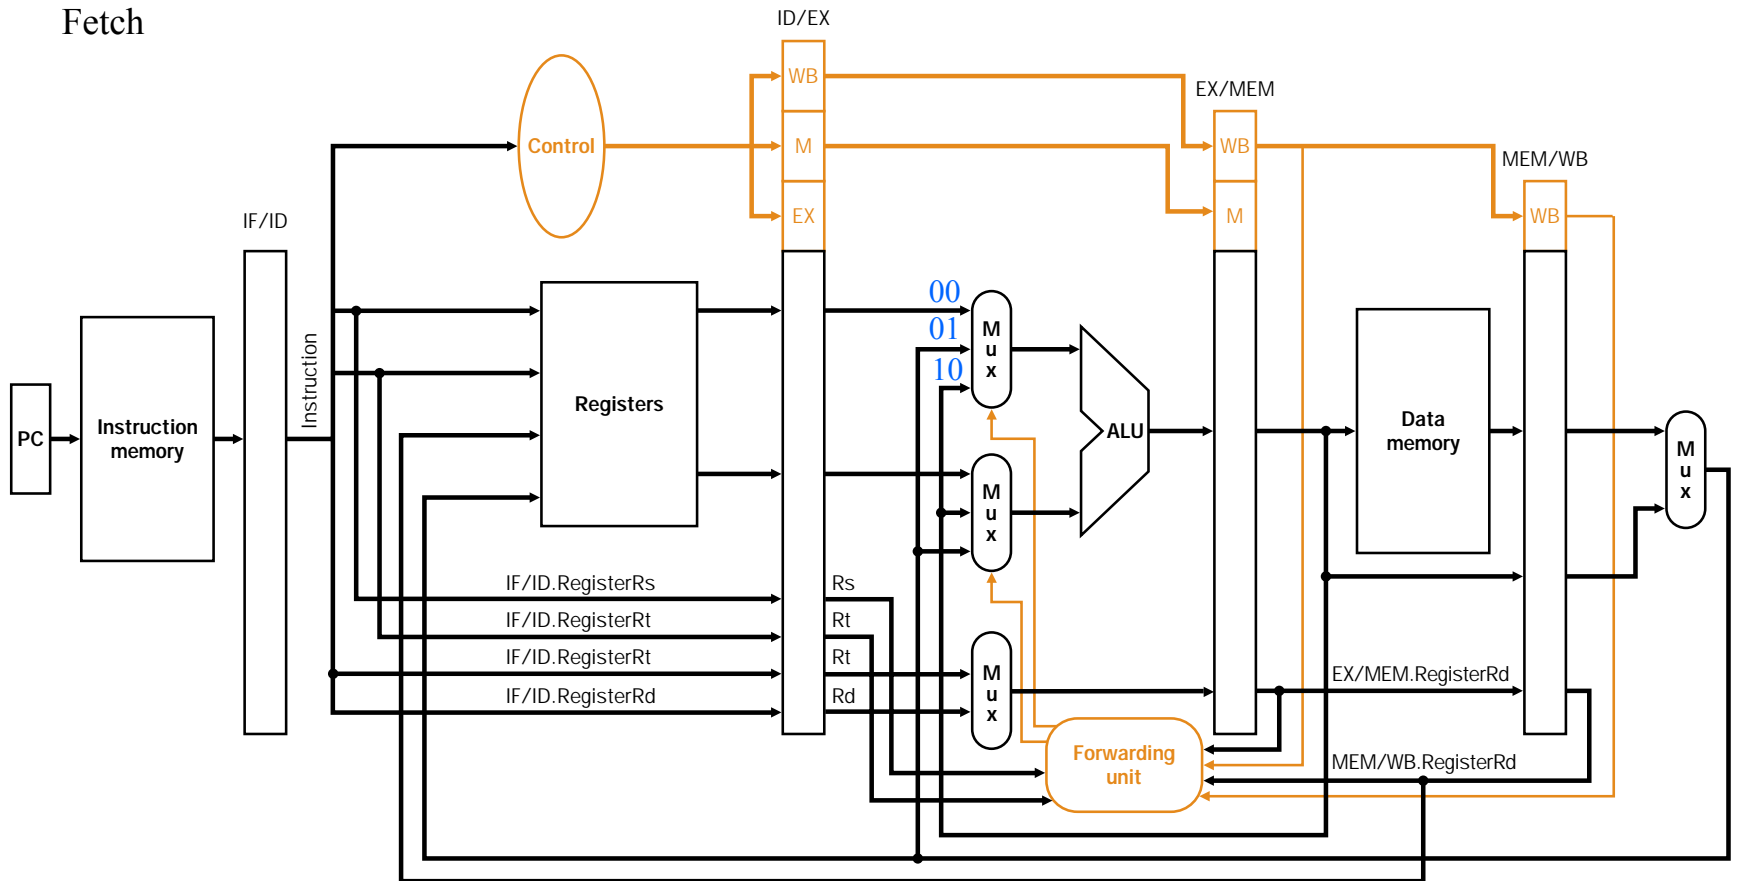
Instruction  
Fetch

sub \$1, \$12, \$3

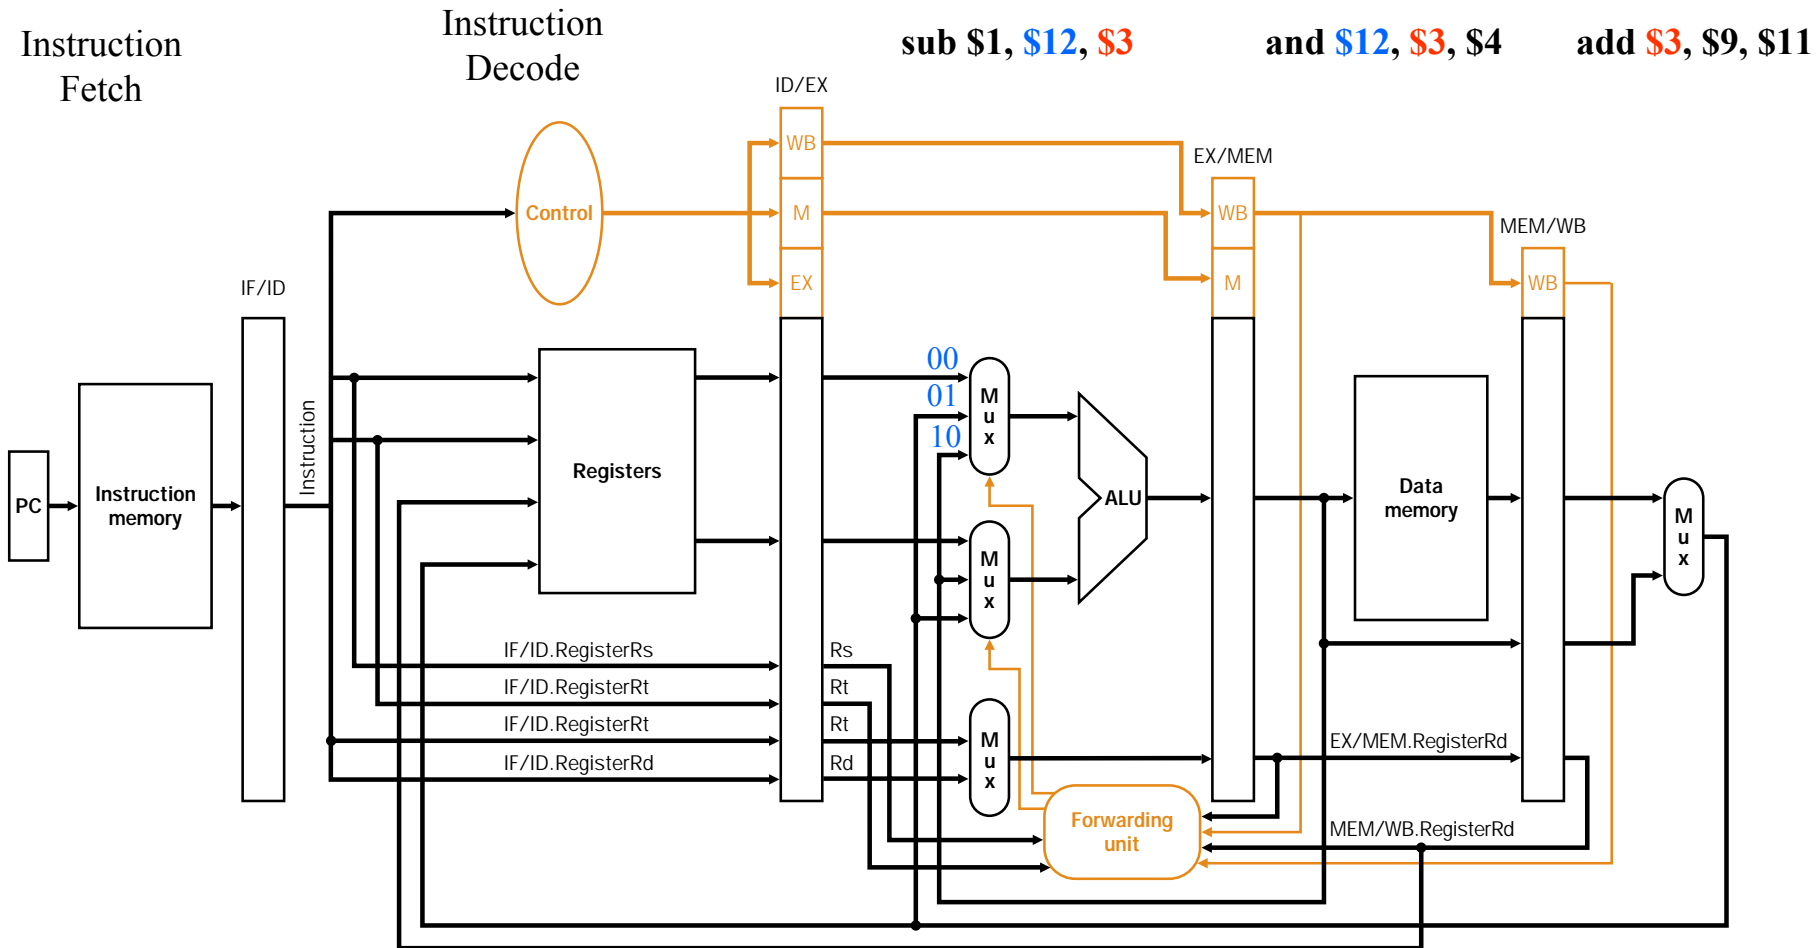
and \$12, \$3, \$4

add \$3, \$8, \$11

Write Back



# Forwarding in Action



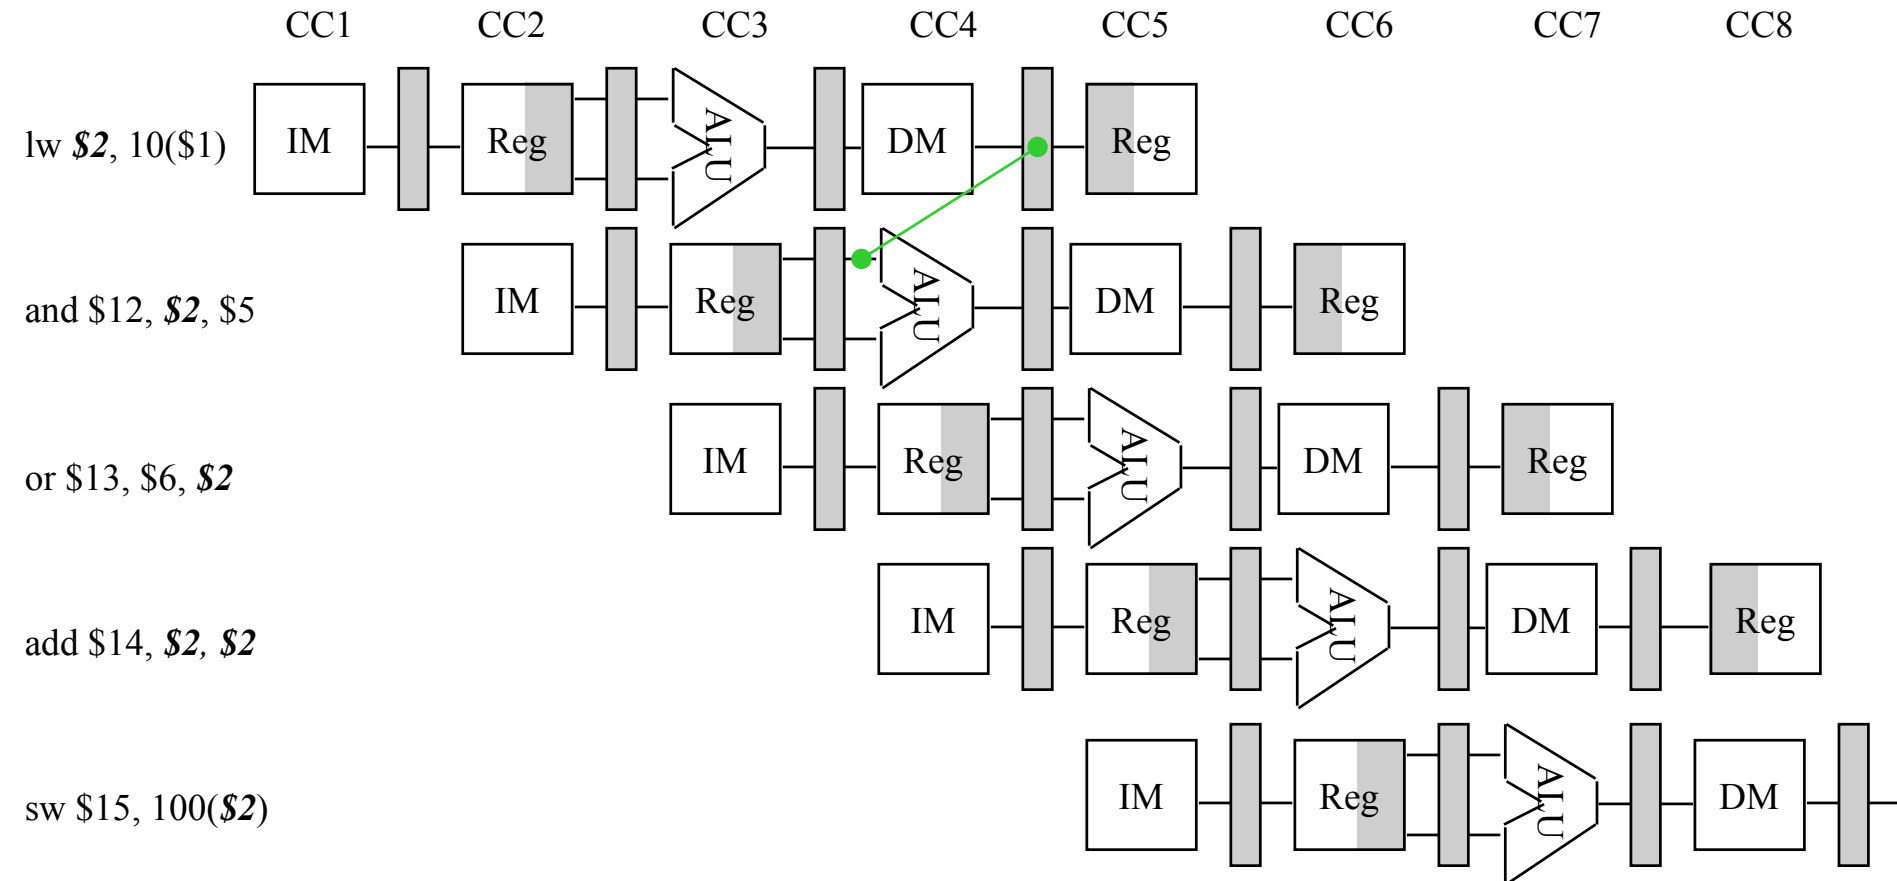


# Forwarding does not eliminate Data Hazard in all cases

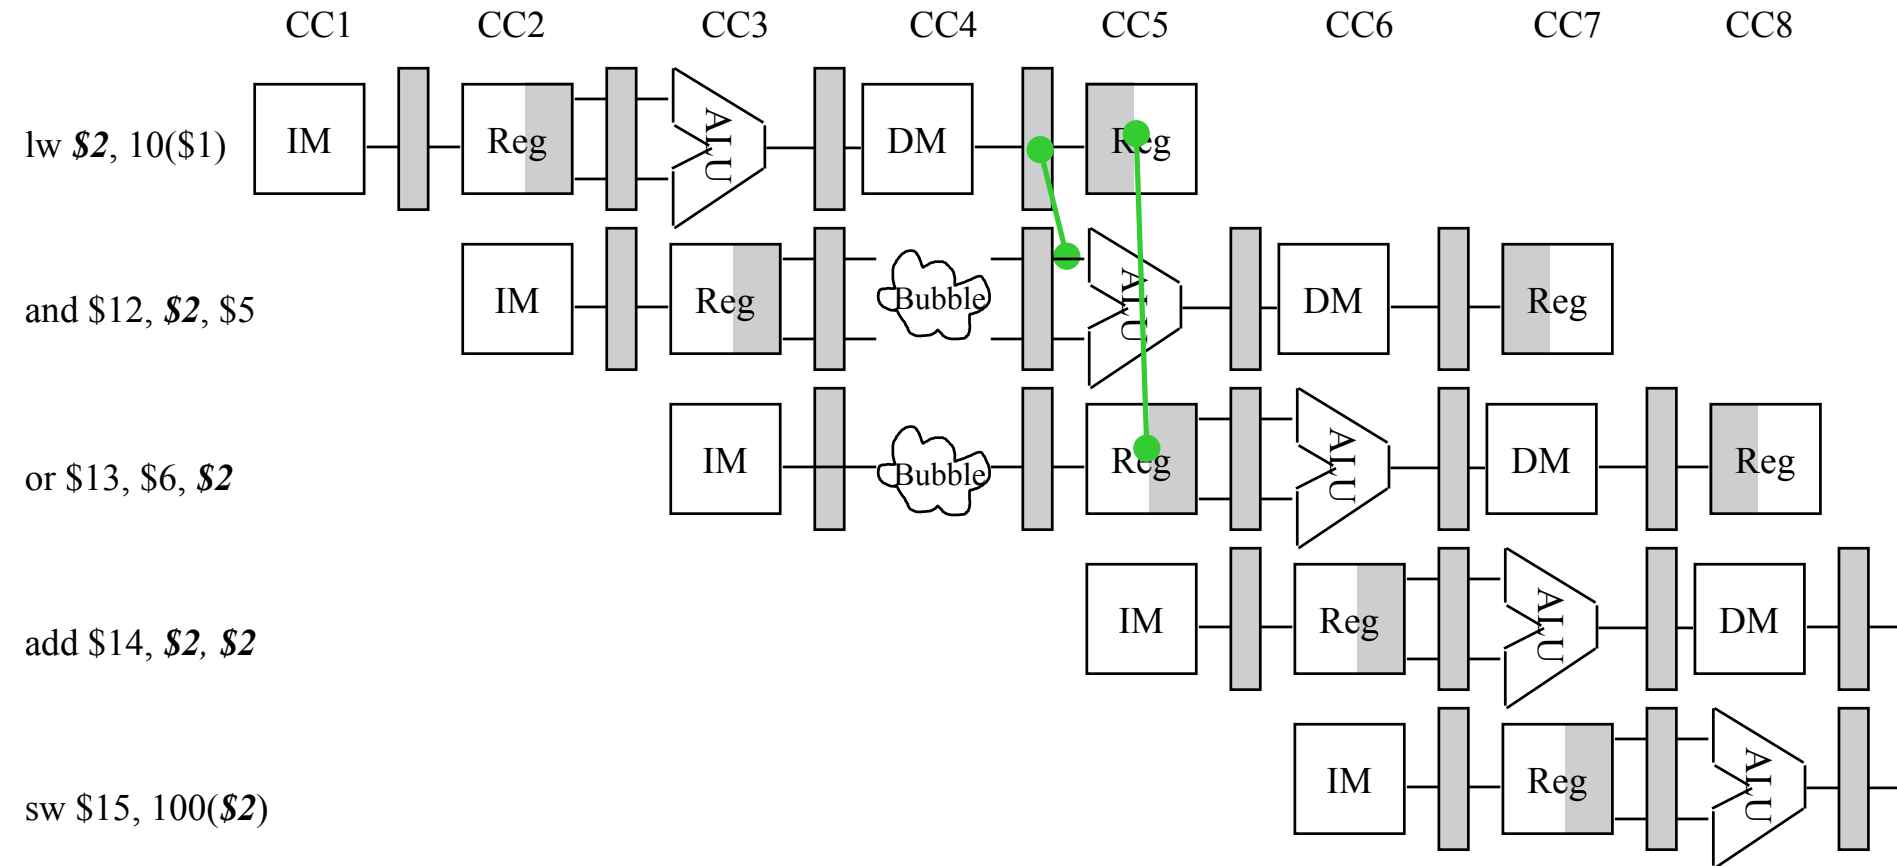
- Consider this code:

```
lw $2, 10($1)
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

# Data Hazard: Load followed by R-type



# Eliminating Data Hazards via Forwarding and stalling

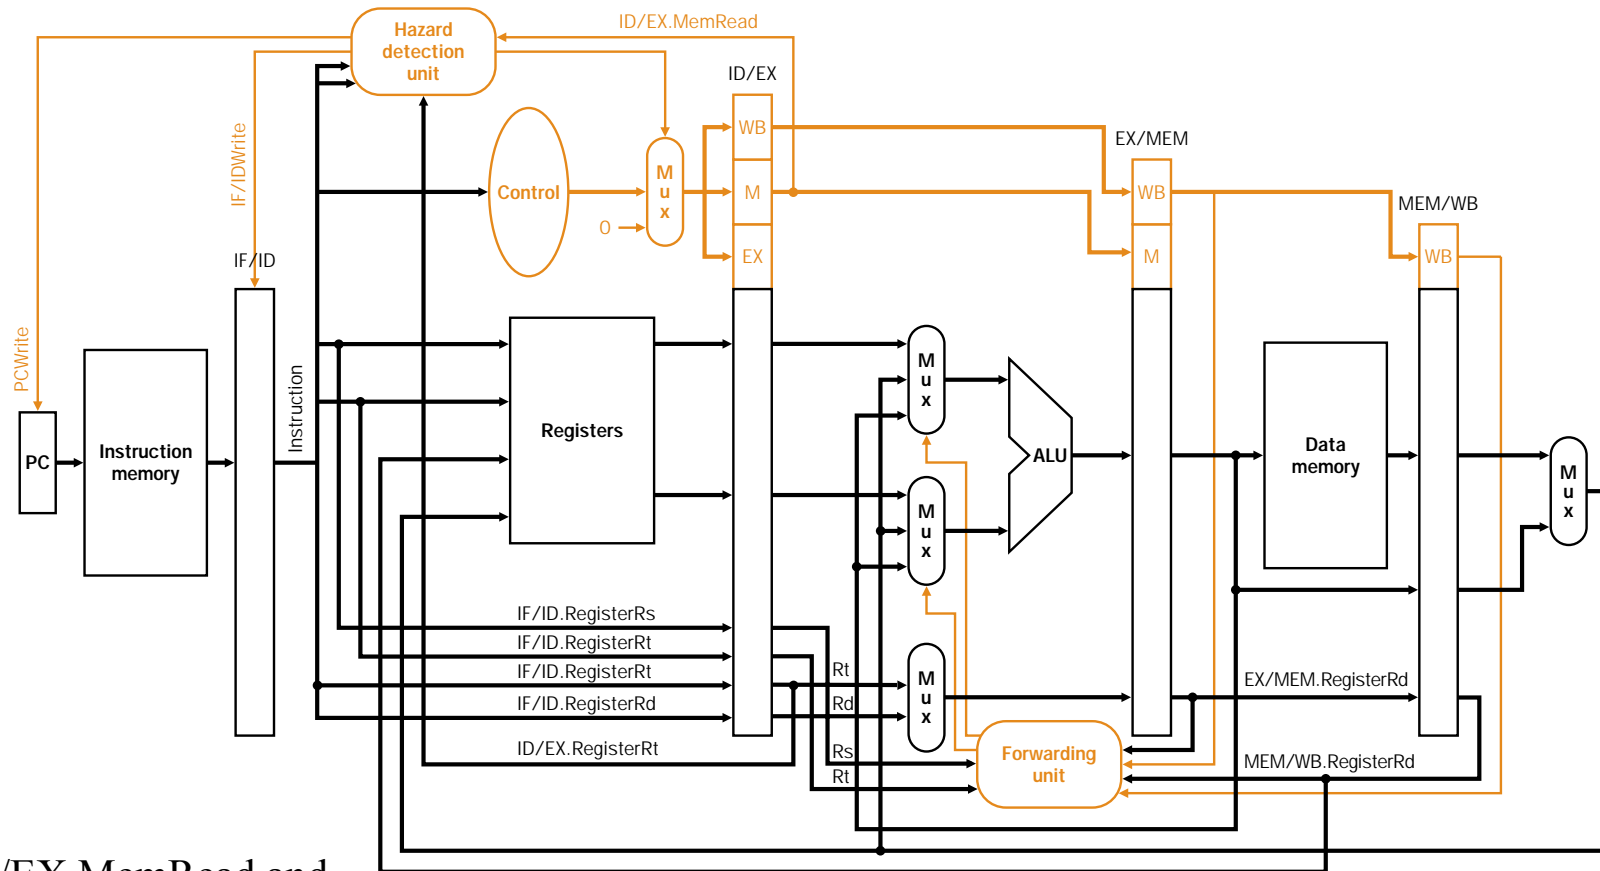


# Pipeline Interlocks

- Not all data hazards can be handled by forwarding
- Pipeline Interlock or Hazard Detection Unit
  - detects a hazard and stalls the pipeline until the hazard is clear
- A stall creates a pipeline bubble:
  - Preventing the IF and ID stages from proceeding
    - don't write the PC ( $PCWrite = 0$ )
    - don't rewrite IF/ID register ( $IF/IDWrite = 0$ )
  - Inserting “nops”
    - set all control signals propagating to EX/MEM/WB to zero (inserts a no-op instruction)

# Hazard Detection Unit

- We can stall the pipeline by keeping an instruction in the same stage




if (ID/EX.MemRead and  
 ((ID/EX.RegisterRt = IF/ID.RegisterRs) or  
 (ID/EX.RegisterRt = IF/ID.RegisterRt)))  
 then stall the pipeline

# Hazard Detection Unit

and \$4, \$2, \$5

lw \$2, 20(\$1)

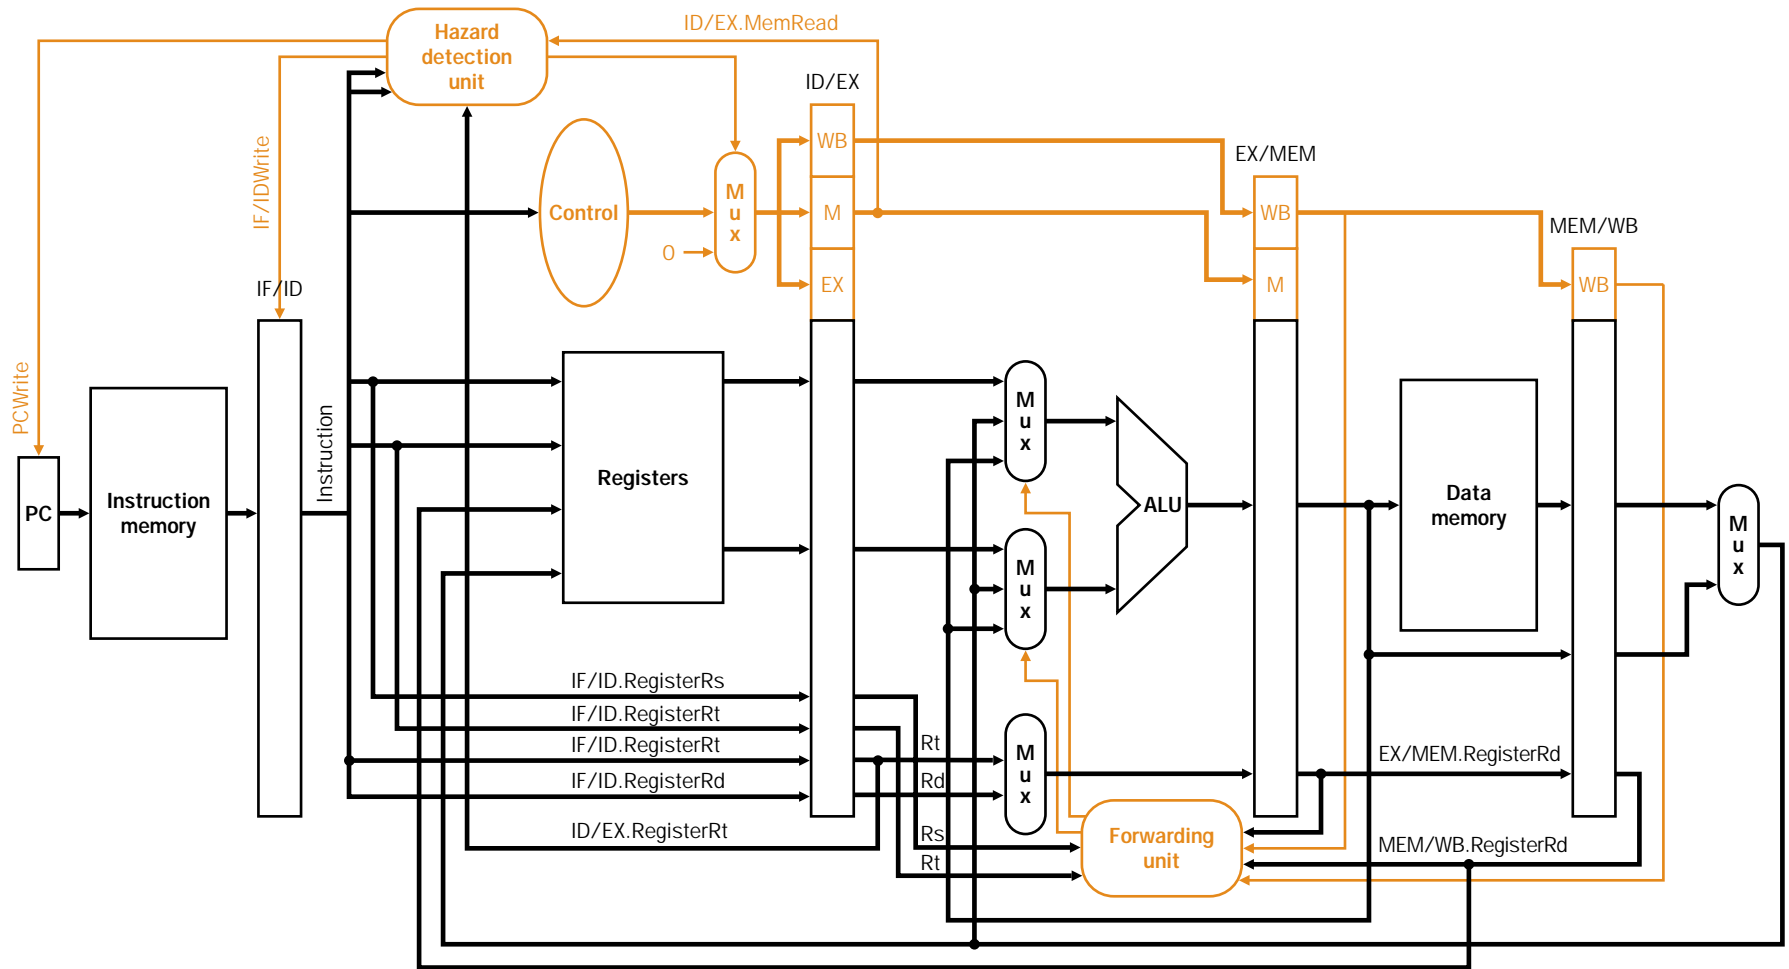


# Hazard Detection Unit

and \$4, \$2, \$5

bubble

lw \$2, 20(\$1)



# Data Hazard Key Points


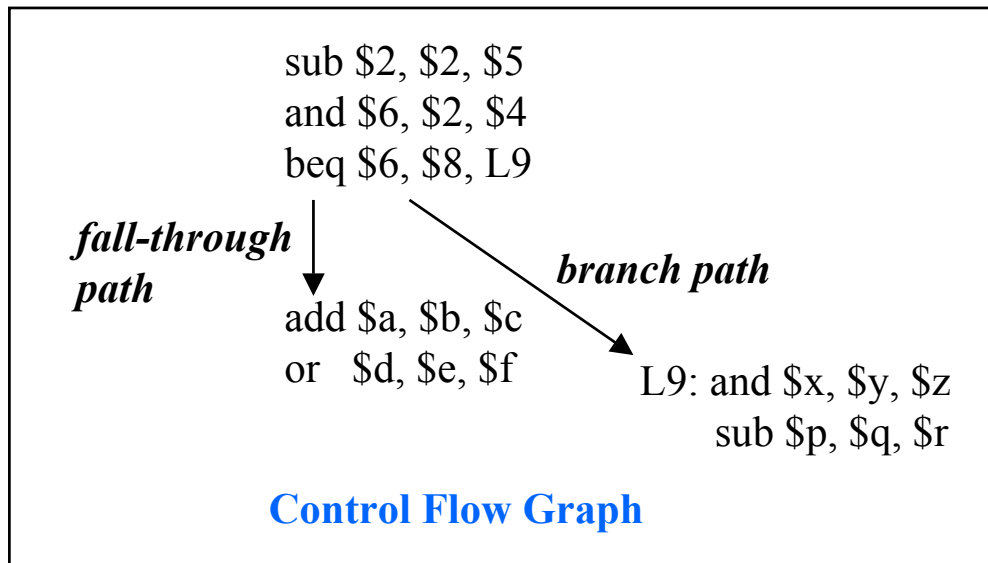
- Pipelining provides high throughput
- Data dependencies cause *data hazards*
- Data hazards can be solved by:
  - Software (nops)
  - Hardware data forwarding
  - Hardware pipeline stalling
- Our processor, and indeed all modern processors, use a combination of forwarding and stalling



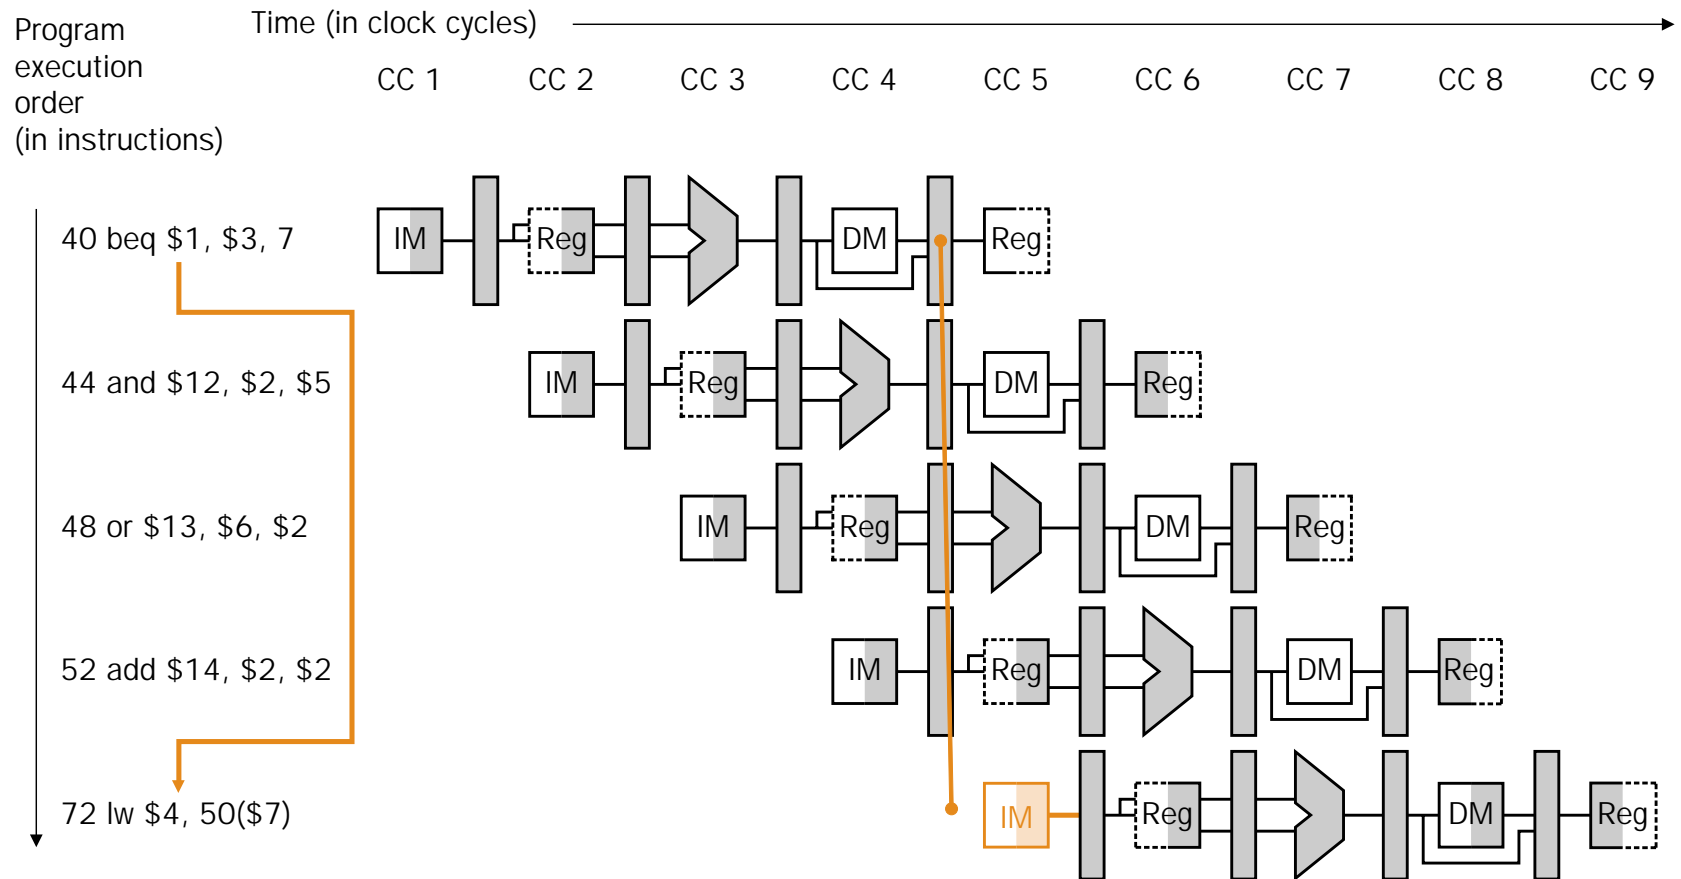
# Control Hazards

# Conditional Branches in a Pipeline

- In a program flow, data computed by certain instructions is used to determine next instruction to execute
  - using conditional branches
- In a pipelined processor, conditional branches result in control hazards



# Impact of a Branch Instruction on the Pipeline

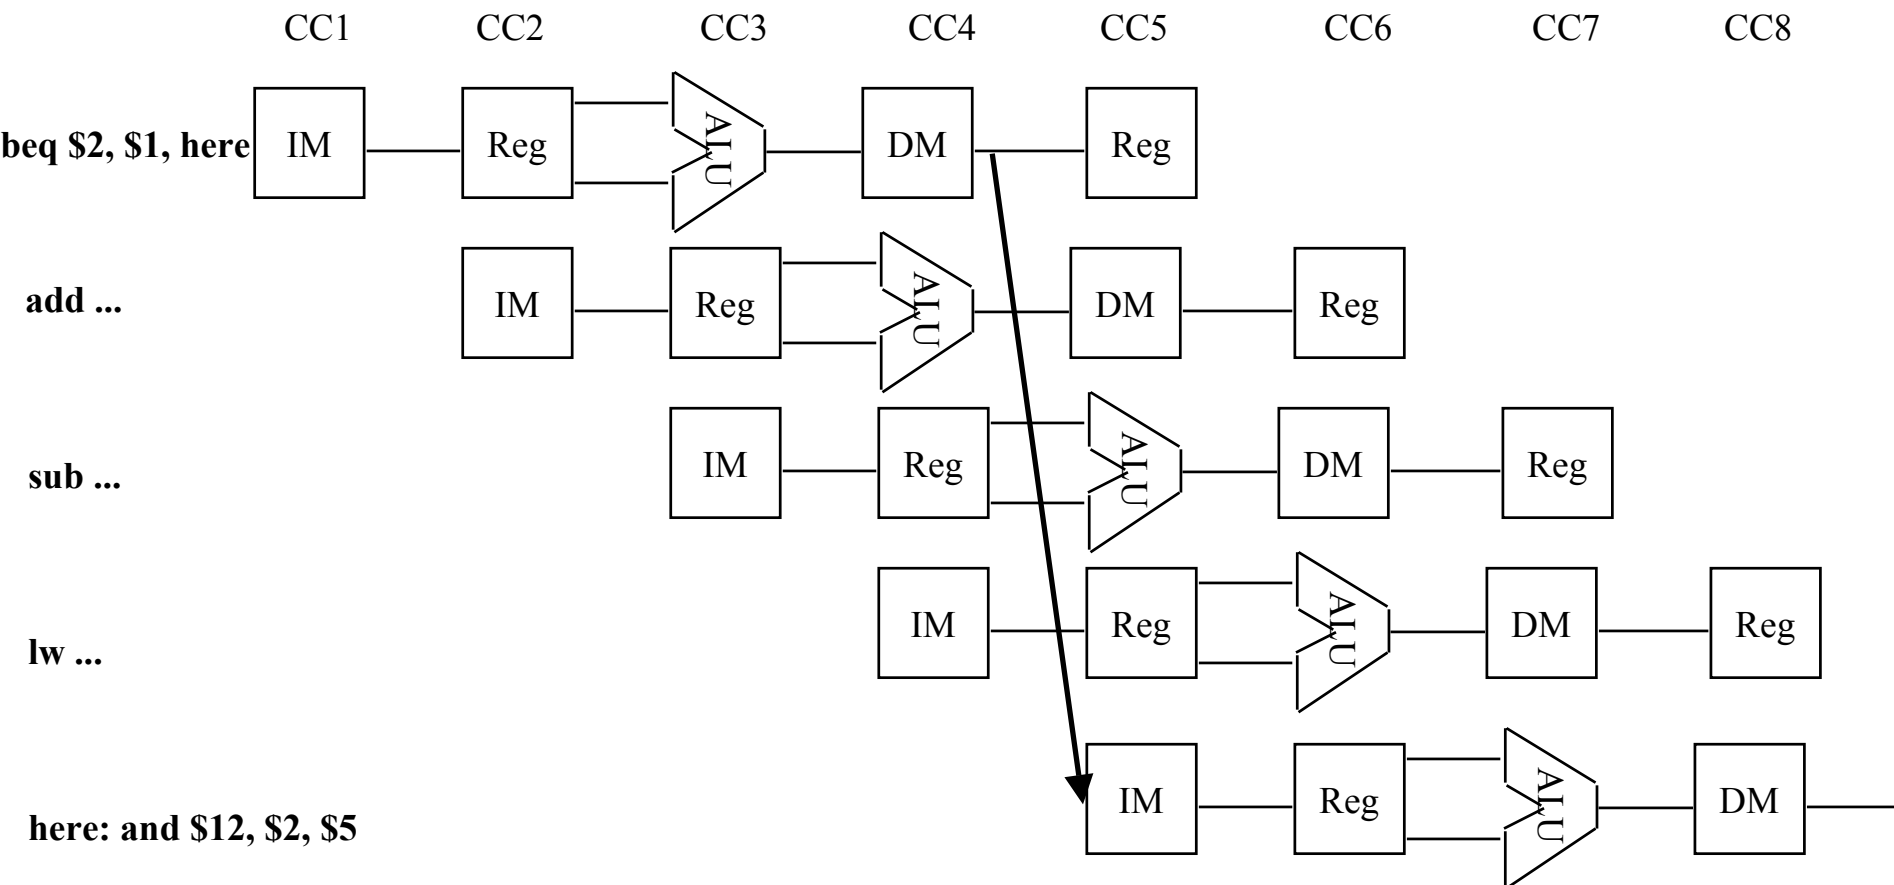


**Decision about whether to branch doesn't occur until the MEM pipeline stage**

# Dealing With Branch Hazards

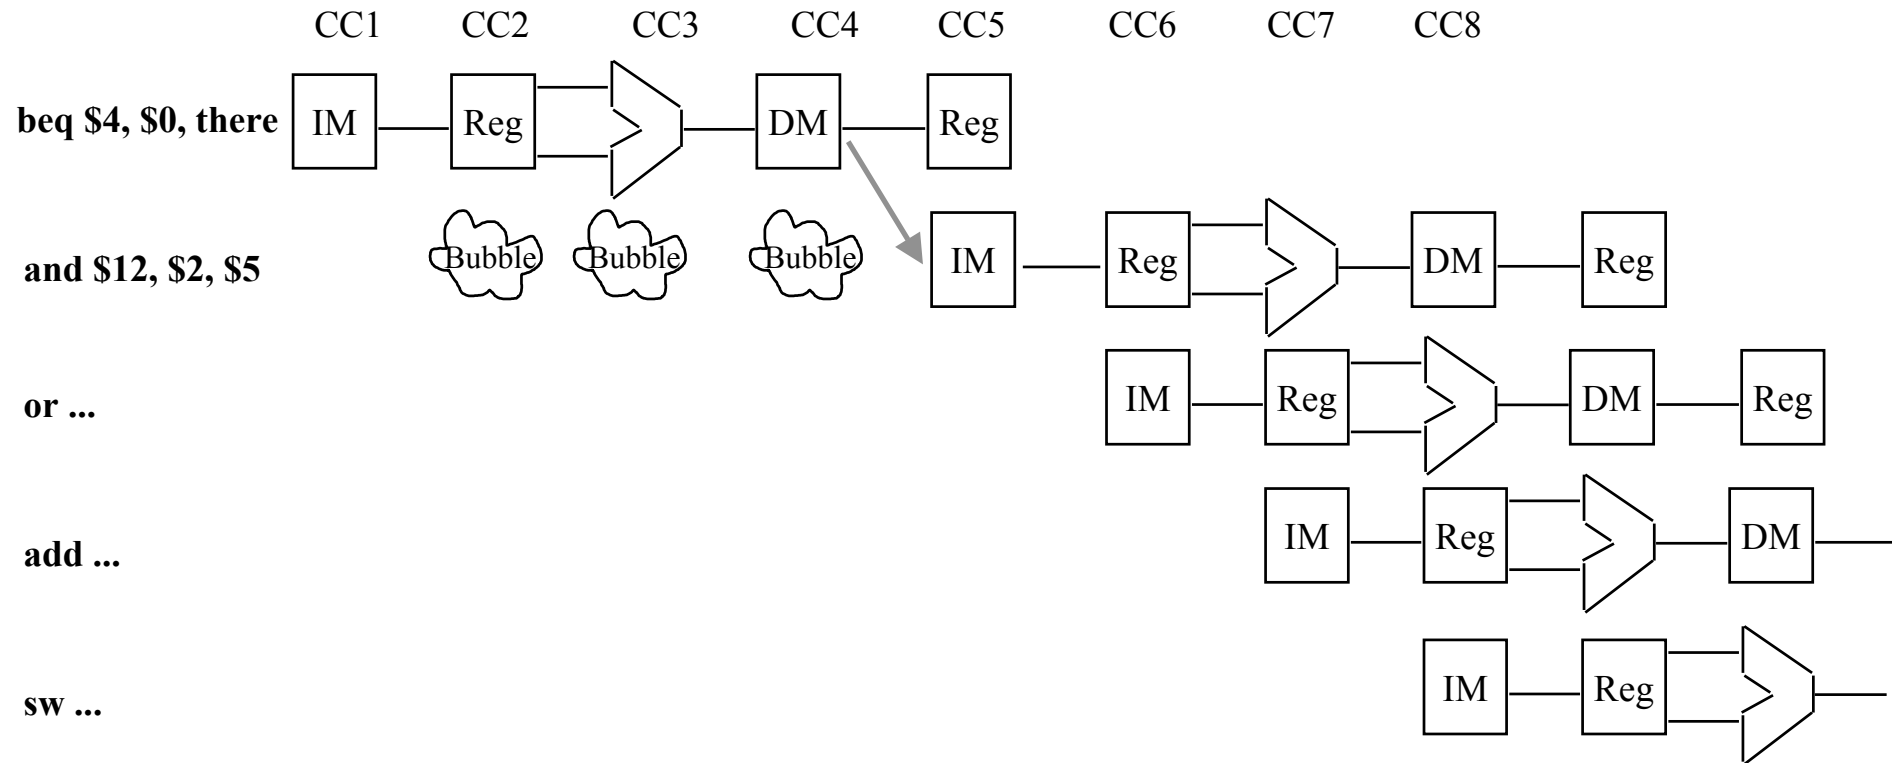
- Software
  - Insert nops,
  - Insert instructions that get executed either way (delayed branch).
- Hardware
  - Stall until you know which direction
    - 3 cycles wasted for every branch
  - Reduce hazard through earlier computation of branch direction
  - Guess which direction
    - assume not taken (easiest)
    - more educated guess based on history (requires that you know it is a branch before it is even decoded!)
  - Ignore the branch for a cycle (branch delay slot)

# Branch Hazards



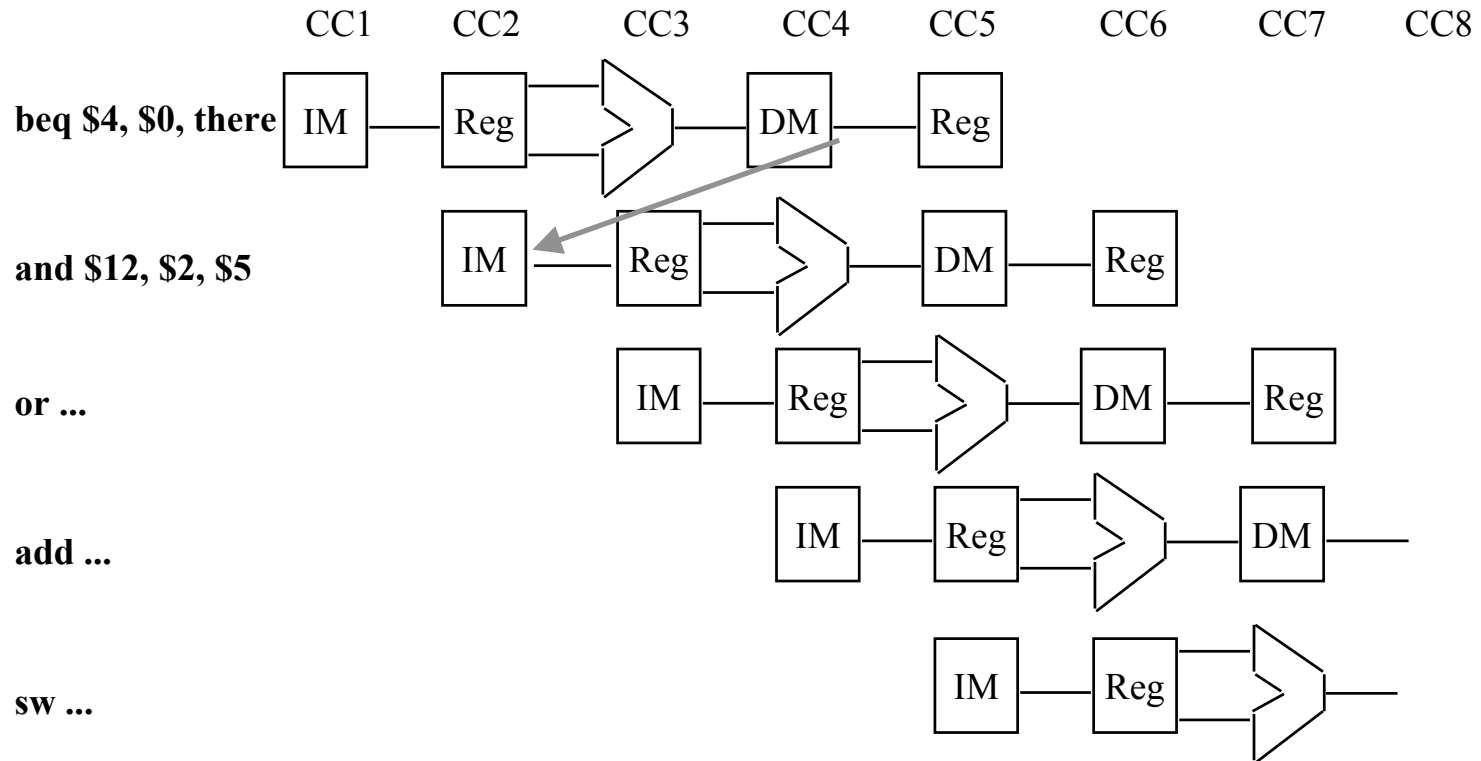
When we decide to branch, other instructions are in the pipeline!

# Stalling for Branch Hazards



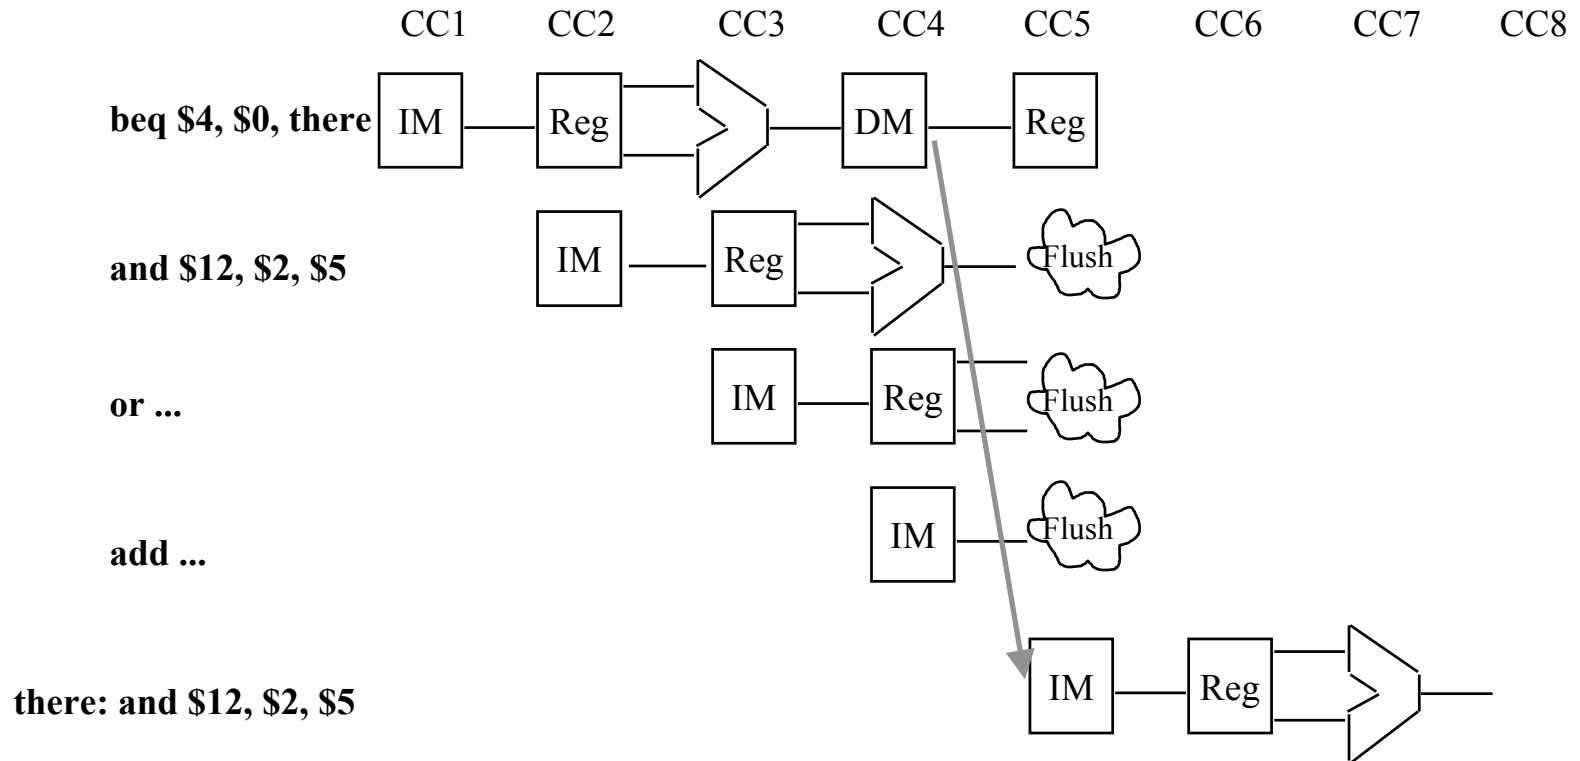
Wastes cycles if branch is not taken

# Assume Branch *Not Taken*



- Works pretty well when you're right

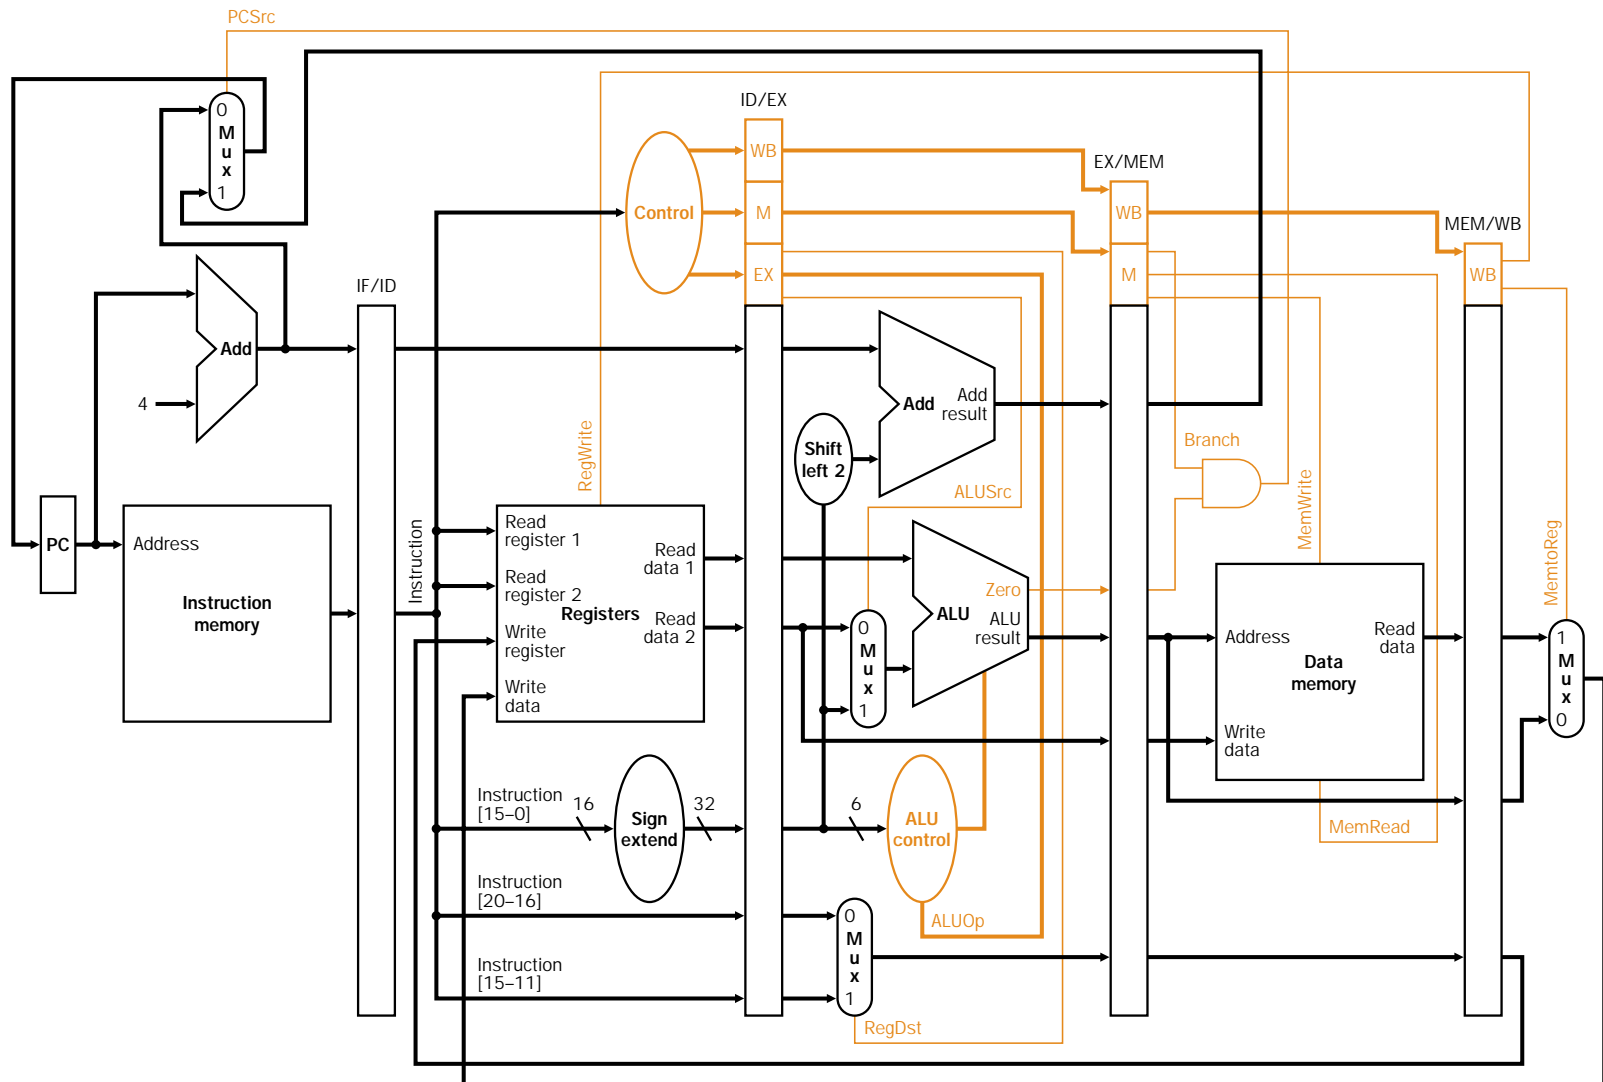
# Assume Branch *Not Taken*



- Same performance as stalling when you're wrong



# Pipelined Datapath and Control



**There is a 3 cycle penalty if a branch is taken**  
**How could we reduce this penalty?**

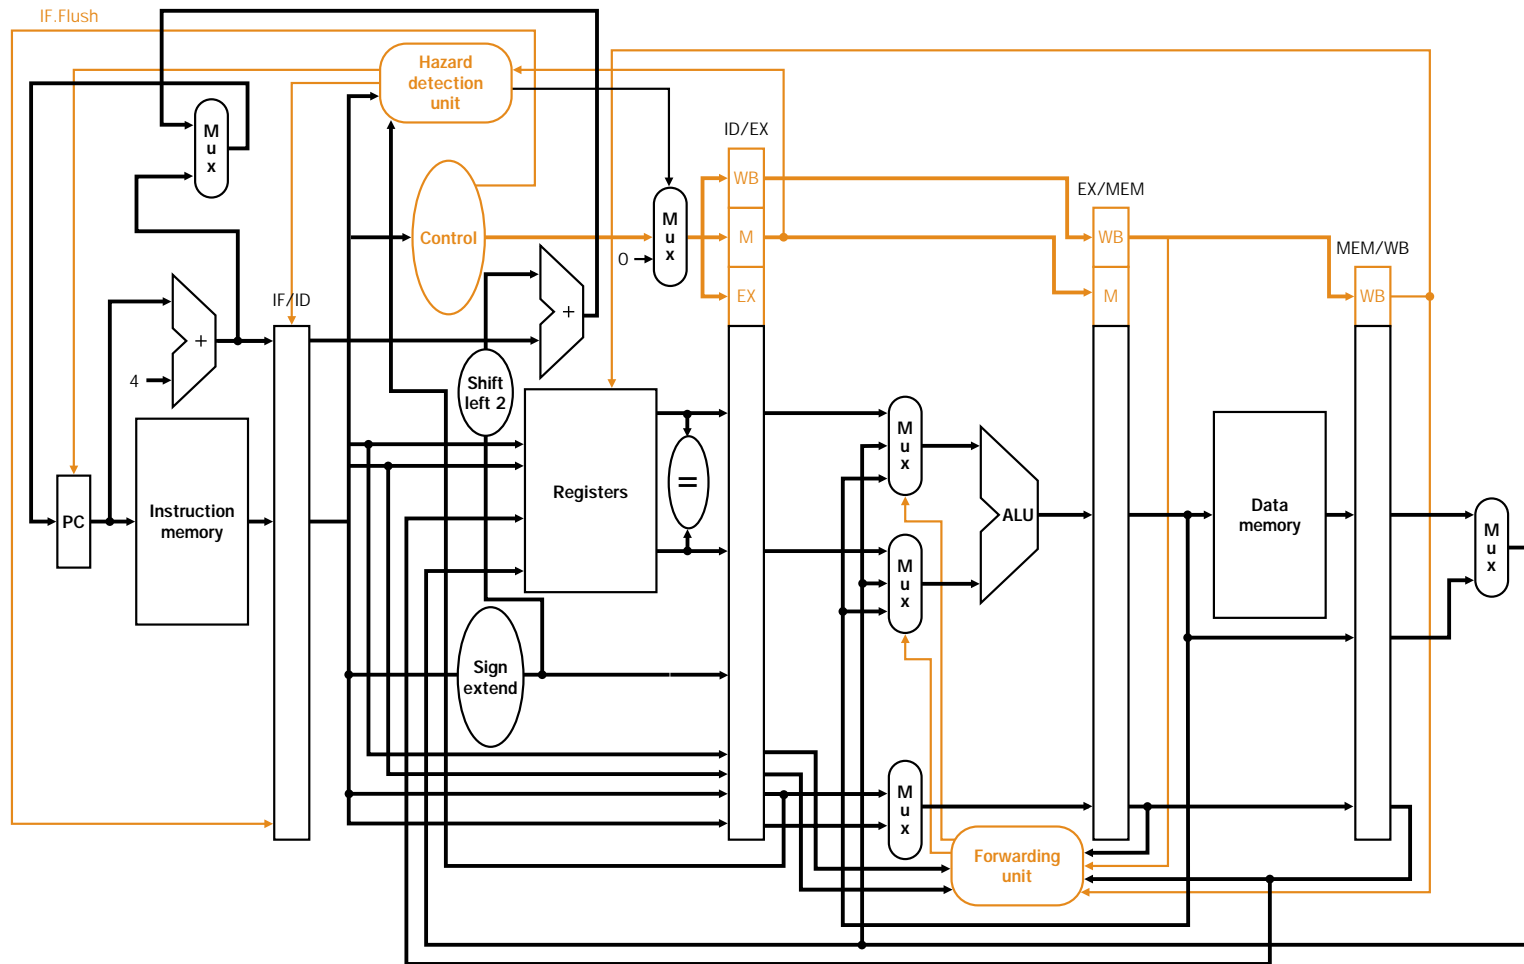
# Reducing the delay of branches

- Resolve the branch in ID stage
  - Move register compare in ID stage
  - Add necessary forwarding muxes and paths
- Implement faster logic to compare registers
  - Current ALU approach
    - Subtract the two registers and check whether the result is zero
    - Slow!
  - Faster approach
    - Exclusive OR the registers and check whether the result is zero
    - Fast, since no carry propagation
- Provide data forwarding
  - Ensure that most recent register values are used in ID stage

# Flush Instructions in the Pipe if a Branch is Taken

- Flushing an instruction means to prevent it from changing any permanent state (registers, memory, PC).
  - Similar to a pipeline bubble...
  - The difference is that we need to be able to insert those bubbles later in the pipeline
- Flushing an instruction on a taken branch
  - Must flush the instruction being fetched in IF stage
  - To flush an instruction,
    - Change all the instruction fields to zero: turn it into nop
    - Let the instruction fields percolate through the pipeline


# Resolving Branch in ID Stage, and Flushing if Branch is Taken



**Note: Forwarding paths and muxes have to be added before registers are compared in ID stage**


## Branch is Taken

36 sub \$10, \$4, \$8  
 40 beq \$1, \$3, 7  
 44 and \$12, \$2, \$5  
 48 or \$13, \$6, \$2  
 52 add \$14, \$2, \$2  
 ...



72 lw \$4, 50(\$7)

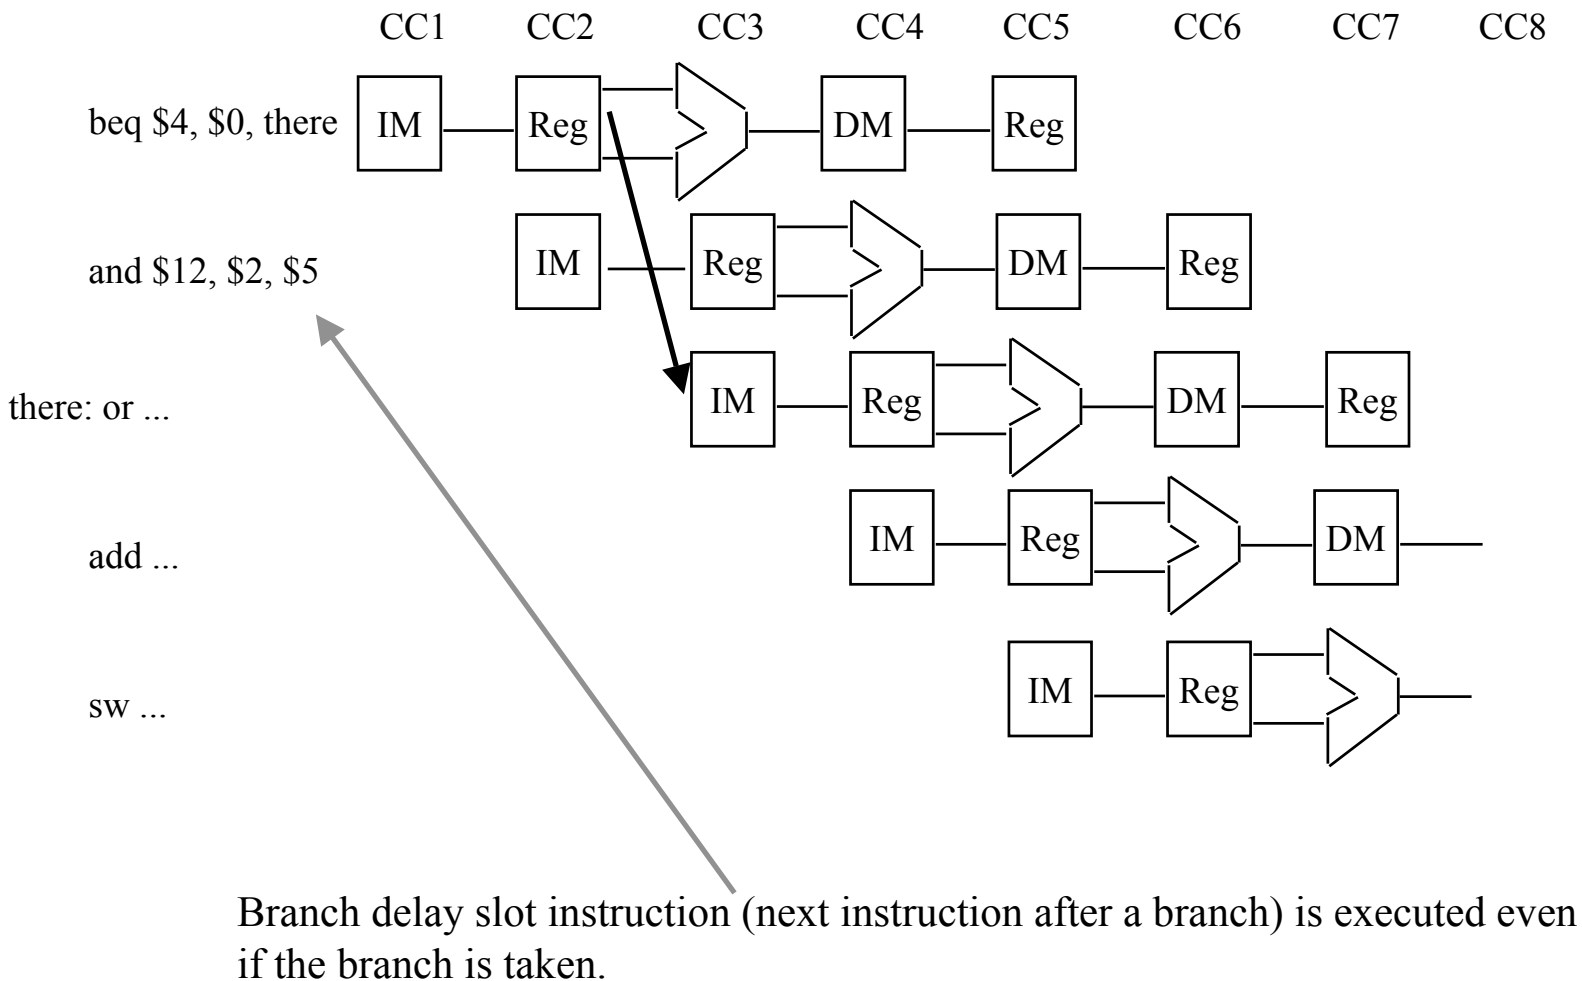
Branch stall reduced  
from 3 cycles to 1 cycle!



# Eliminating the 1 Cycle Branch Stall

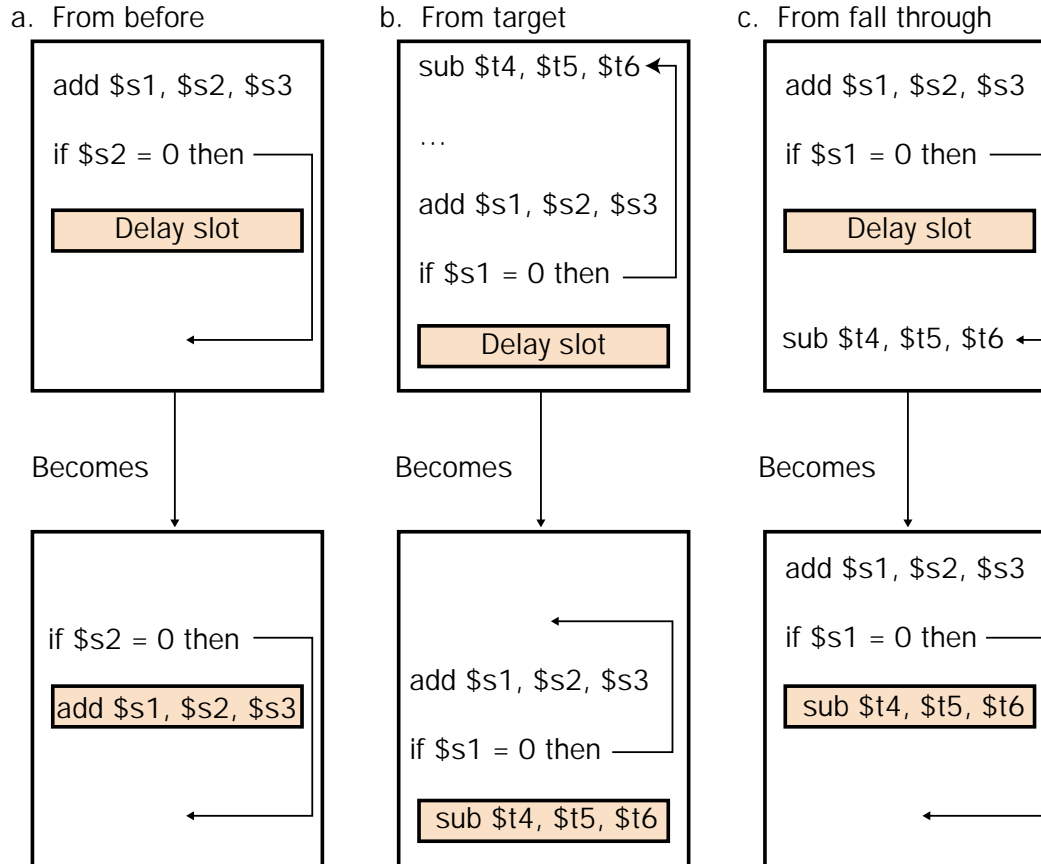
- There's no rule that says we have to see the effect of the branch immediately. Why not wait an extra instruction before branching?
- The original SPARC and MIPS processors each used a single *branch delay slot* to eliminate single-cycle stalls after branches.
- The instruction after a conditional branch is always executed in those machines, regardless of whether the branch is taken or not!
- This works great for this implementation of the architecture, but becomes a permanent part of the ISA.
- What about the MIPS R10000, which has a 5-cycle branch penalty, and executes 4 instructions per cycle?

# Branch Delay Slot



# Scheduling Branch Delay Slot

The branch delay slot is only useful if you can find something to put there. If you can't find anything, you must put a *nop* to insure correctness.



For b and c, \$t4 must be an unused temporary register



# Importance of Branch Prediction

- 15 to 20% of all instructions are branches
- MIPS
  - branch stall of 1 cycle, 1 instruction issued per cycle
  - delayed branch
- Recent processors
  - 3-4 cycle hazard, 1-2 instructions issued per cycle
  - cost of branch misprediction goes up
- Pentium Pro
  - 12+ cycle misprediction penalty, 3 instructions issued per cycle
  - HUGE penalty for mispredicting a branch
  - 36+ issue slots wasted

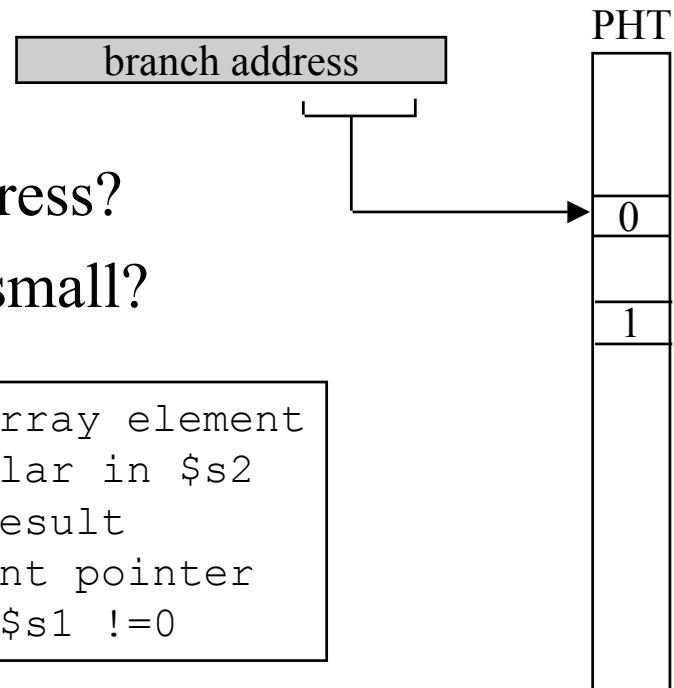
# Predicting Branch Direction

- Easiest
  - always not taken, always taken
  - forward not taken, backward always taken
    - Appropriate for loops
  - compiler predicted (branch likely, branch not likely)
- Next easiest
  - Record 1-bit history of whether the branch was taken or not
    - 1-bit predictor
    - For a loop, the predictor is incorrect twice

# 1-bit Pattern History Table

- Uses low bits of branch address to choose an entry
- The entry has 1 branch prediction bit
- Size is small, e.g. 1 bits by N (e.g. 4K)

- Why not use all bits of branch address?
- What happens when the table too small?

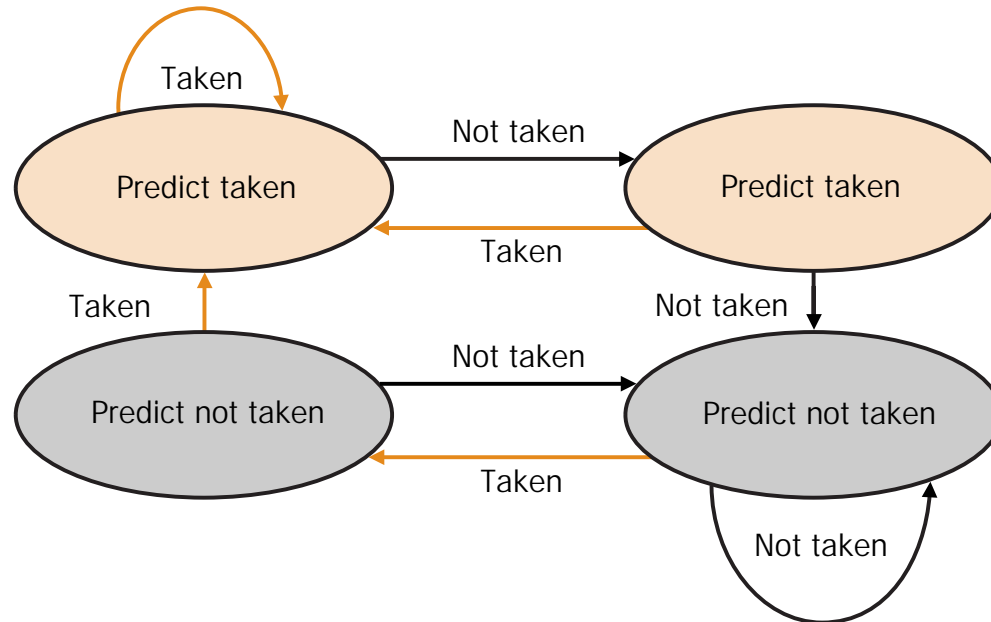


```
Loop:lw    $t0, 0($s1)    # $t0 = array element
      addu $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)    # store result
      addi $s1, $s1, -4    # decrement pointer
      bne   $s1, $zero, Loop # branch $s1 !=0
```

- Prediction is incorrect twice

# 2-bit Branch Prediction Scheme

Branch prediction has to be incorrect twice before it is changed



# Control Hazards -- Key Points


- Control (or branch) hazards arise because we must fetch the next instruction before we know if we are branching or where we are branching.
- Control hazards are detected in hardware.
- We can reduce the impact of control hazards through:
  - early detection of branch address and condition
  - branch prediction
  - branch delay slots

# Exceptions

# Exception Handling in the Pipeline

- Consider arithmetic overflow exception
  - add \$2, \$6, \$5
- Extra hardware
  - Note: add is in EX stage
  - Flush instructions that follow add
    - In IF stage, assert IF.flush
    - In ID stage, use mux added for stall (OR ID.flush)
    - In EX stage, use EX.flush signal
  - Transfer control to PC = 0x4000 0040
  - Save PC + 4 in EPC
  - Save exception cause in Cause Register

# Datapath and Control for Exceptions





# Exception Handling in a Pipeline


```

0x40 sub $11, $2, $4
0x44 and $12, $2, $5
0x48 or $13, $2, $6
0x4c add $1, $2, $1
0x50 slt $15, $6, $7
0x54 lw $16, 50($7)
  
```

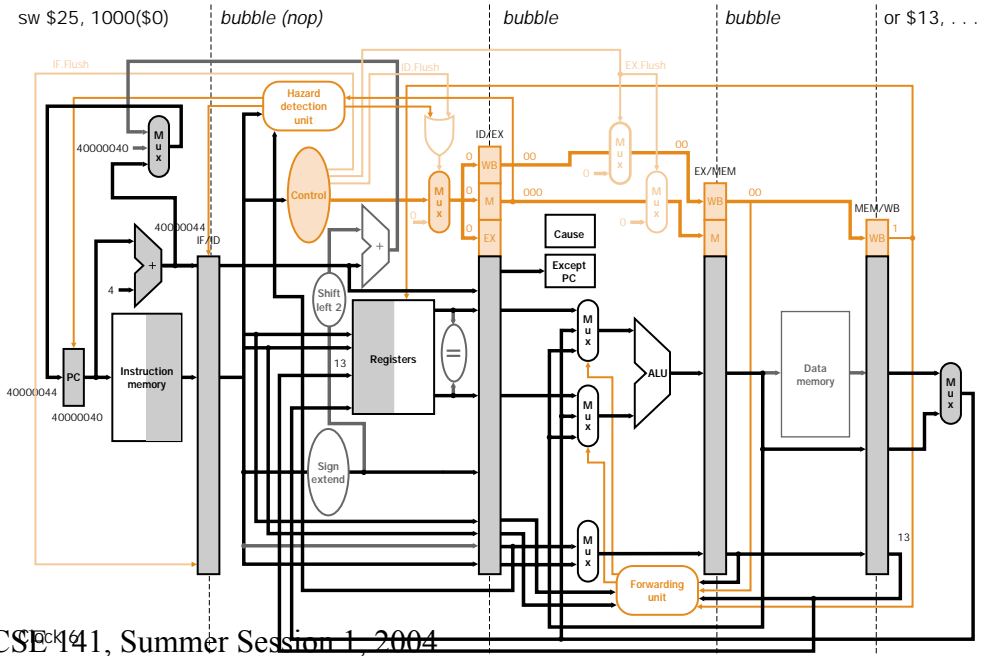
```

0x40000040 sw $25, 1000($0)
0x40000044 sw $26, 1004($0)
  
```

**Note: ALU overflow signal  
is input the control unit**



Clock 5



Clock 6

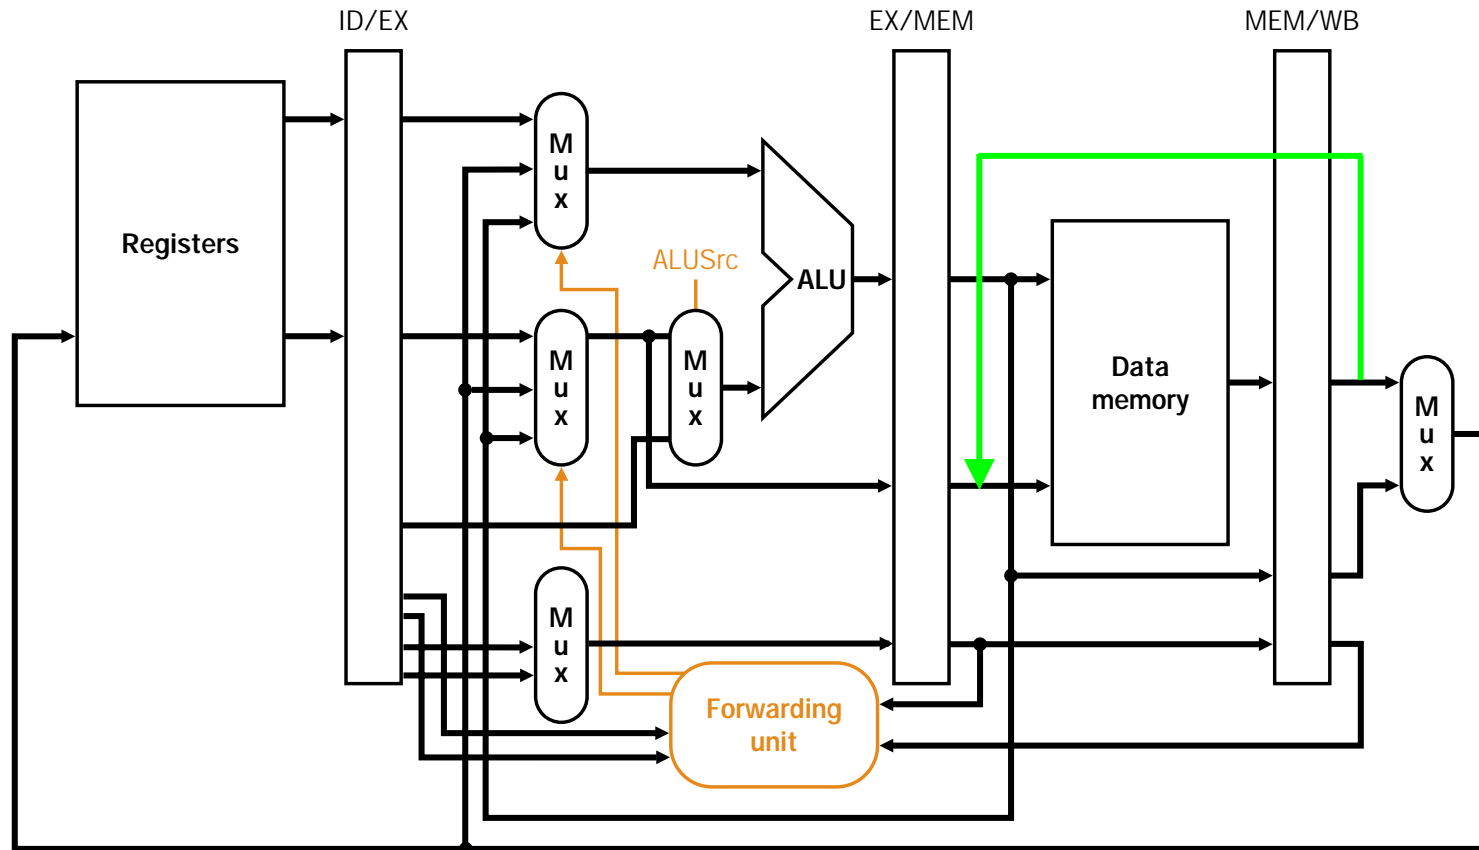
# Issues in Handling an Exception

- Five instructions are active in the pipeline
- Multiple exceptions may occur
  - Earliest instruction is generated interrupted
- Exceptions are detected in different stage of the pipeline
  - Undefined instruction is discovered in ID stage
  - Overflow is detected in EX stage
  - Kernel call (i.e. OS call) is detected in EX stage
- Precise exception
  - EPC saves PC of the instruction that caused exception
  - This is required for virtual memory
- Imprecise exception
  - EPC may not save PC of the instruction that caused exception
    - For ease of implementation

# Summary: Solutions to Hazards

- Control Hazards
  - Stall on branch
  - Predict branch
- Data Hazard
  - Operand forwarding or bypassing

# $M \Rightarrow M$ Forwarding for $LW \Rightarrow SW$ (Exercise 6.20 in the Textbook)



# Example

	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	C11	C12	C13	C14	C15
<b>ADD R1, R2, R3</b>															
<b>SW R1, 1000(R2)</b>															
<b>LW R7, 2000(R2)</b>															
<b>ADD R5, R7, R1</b>															
<b>LW R8, 2004(R2)</b>															
<b>SW R7, 2008(R8)</b>															
<b>ADD R8, R8, R2</b>															
<b>LW R9, 1012(R8)</b>															
<b>SW R9, 1016(R8)</b>															
Pramod Argade															